

Recommendation Systems

based on slides by Alex Smola
Yahoo! Research and ANU

Significant content courtesy of Yehuda Koren

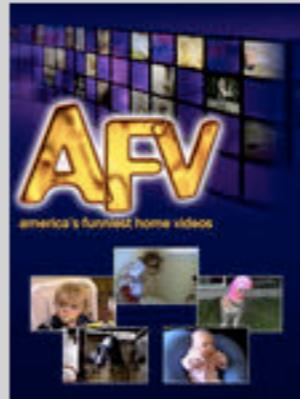
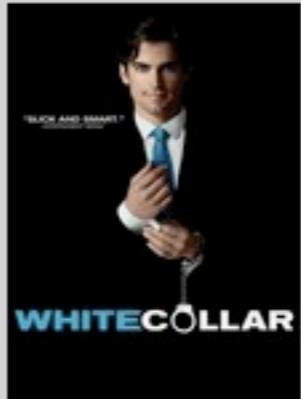
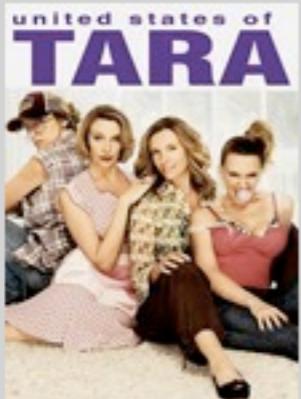
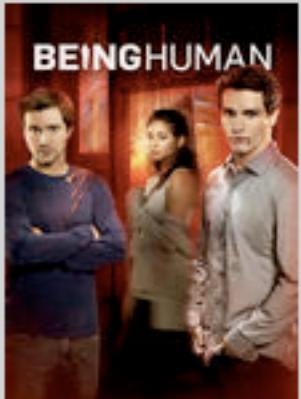
Outline

- Neighborhood methods
 - User / movie similarity
 - Iteration on graph
- Matrix Factorization
 - Singular value decomposition
 - Convex reformulation
- Improvements

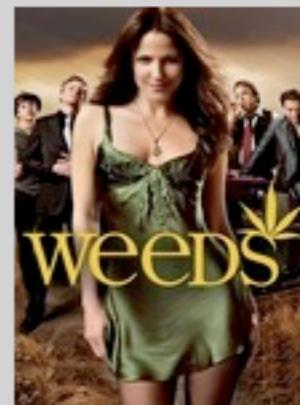
Why

Thousands of movies and TV episodes including these:

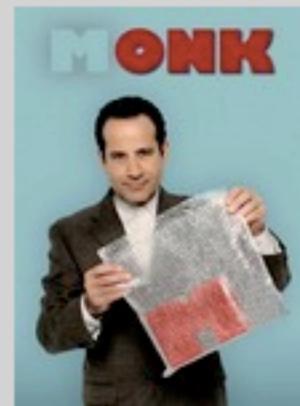
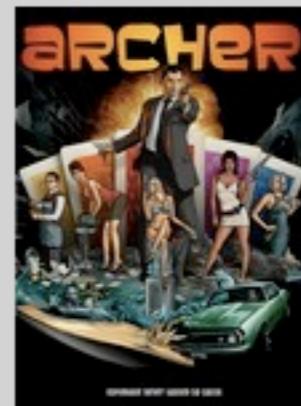
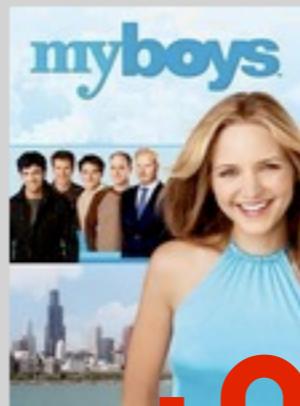
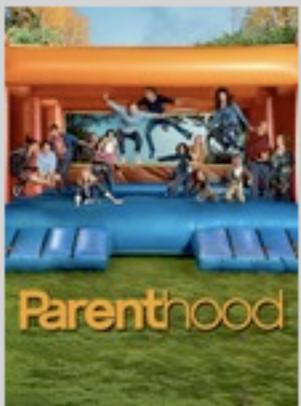
New Arrivals in TV



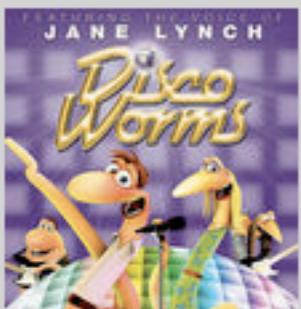
TV Drama



TV Comedy



Children & Family



Netflix

Shop All Departments

Search

Amazon Instant Video

GO



Cart

Wish List

Amazon Instant Video

Most Popular

Getting Started

Watch Anywhere

Prime Instant Videos

Your Video Library

Passes and Pre-orders

Get Help

DVD & Blu-ray



Three Kings

(398 customer reviews)

George Clooney, Mark Wahlberg, Ice Cube conspire to steal a huge cache of gold hidden near their desert base.



Starring: George Clooney, Mark Wahlberg

Directed by: David O. Russell

Runtime: 1 hour 56 minutes

Release year: 1999

Studio: Warner Bros.



Also available in **HD** with [Amazon Instant Video on Your TV](#)



Your Amazon Prime membership now includes unlimited, commercial-free, instant streaming of thousands of [movies](#) and [TV shows](#) at no additional cost.

amazon
instant video

Prime instant videos

Watch now **\$0.00**
unlimited streaming

48 hour rental

1-Click® \$2.99

Buy movie

1-Click® \$9.99

[Learn more about renting and buying](#)

Add to Wish List

[Send us Feedback]

Have a promotion code? [Redeem a gift card or promotion code](#) [View Balance](#)

Customers Who Bought This Item Also Bought



[Tower Heist](#) Amazon Instant Video ~ Eddie Murphy
 (42)
\$3.99



[Syriana](#) Amazon Instant Video ~ George Clooney
 (355)
\$2.99



[Five Minutes of Heaven](#) Amazon Instant Video ~ Liam Neeson
 (70)
\$2.99



[Foolproof](#) Amazon Instant Video ~ Ryan Reynolds
 (81)
\$2.99



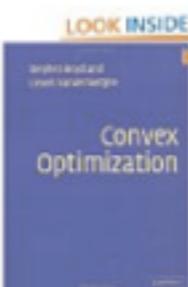
[The Recruit](#) Amazon Instant Video ~ Al Pacino
 (161)
\$1.99



Page 1 of 19

Just For Today[Browse Recommended](#)**Recommendations**[Amazon Instant Video](#)[Appstore for Android](#)[Baby](#)[Beauty](#)[Books](#)[Books on Kindle](#)[Camera & Photo](#)[Clothing & Accessories](#)[Computers & Accessories](#)[Electronics](#)[Grocery & Gourmet Food](#)[Health & Personal Care](#)[Home Improvement](#)[Industrial & Scientific](#)[Jewelry](#)[Kitchen & Dining](#)[MP3 Downloads](#)[Magazine Subscriptions](#)[Movies & TV](#)[Music](#)[Musical Instruments](#)[Office & School Supplies](#)[Patio, Lawn & Garden](#)[Shoes](#)[Software](#)[Sports & Outdoors](#)[Toys & Games](#)[Video Games](#)[Watches](#)These recommendations are based on [items you own](#) and more.view: [All](#) | [New Releases](#) | [Coming Soon](#)[More results](#)

1.

**Convex Optimization**

by Stephen Boyd (March 8, 2004)

Average Customer Review: (13)

In Stock

List Price: \$84.00**Price:** \$68.13[44 used & new from \\$61.32](#)[Add to Cart](#)[Add to Wish List](#) I own it Not interested Rate this itemRecommended because you purchased [Nonlinear Programming](#) and more ([Fix this](#))

2.

**Probabilistic Graphical Models: Principles and Techniques (Adaptive Computation and Machine Learning series)**

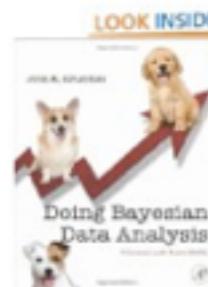
by Nir Friedman (July 31, 2009)

Average Customer Review: (11)

In Stock

List Price: \$95.00**Price:** \$93.55[47 used & new from \\$91.33](#)[Add to Cart](#)[Add to Wish List](#) I own it Not interested Rate this itemRecommended because you purchased [Nonlinear Programming](#) and more ([Fix this](#))

3.

**Doing Bayesian Data Analysis: A Tutorial with R and BUGS**

by John K. Kruschke (November 10, 2010)

Average Customer Review: (15)

In Stock

List Price: \$89.95**Price:** \$77.98[44 used & new from \\$68.40](#)[Add to Cart](#)[Add to Wish List](#) I own it Not interested Rate this itemRecommended because you purchased [Bayesian Nonparametrics](#) and more ([Fix this](#))

4.

**Parallel and Distributed Computation: Numerical Methods (Optimization and Neural Computation)**

by Dimitri P. Bertsekas (January 1, 1997)

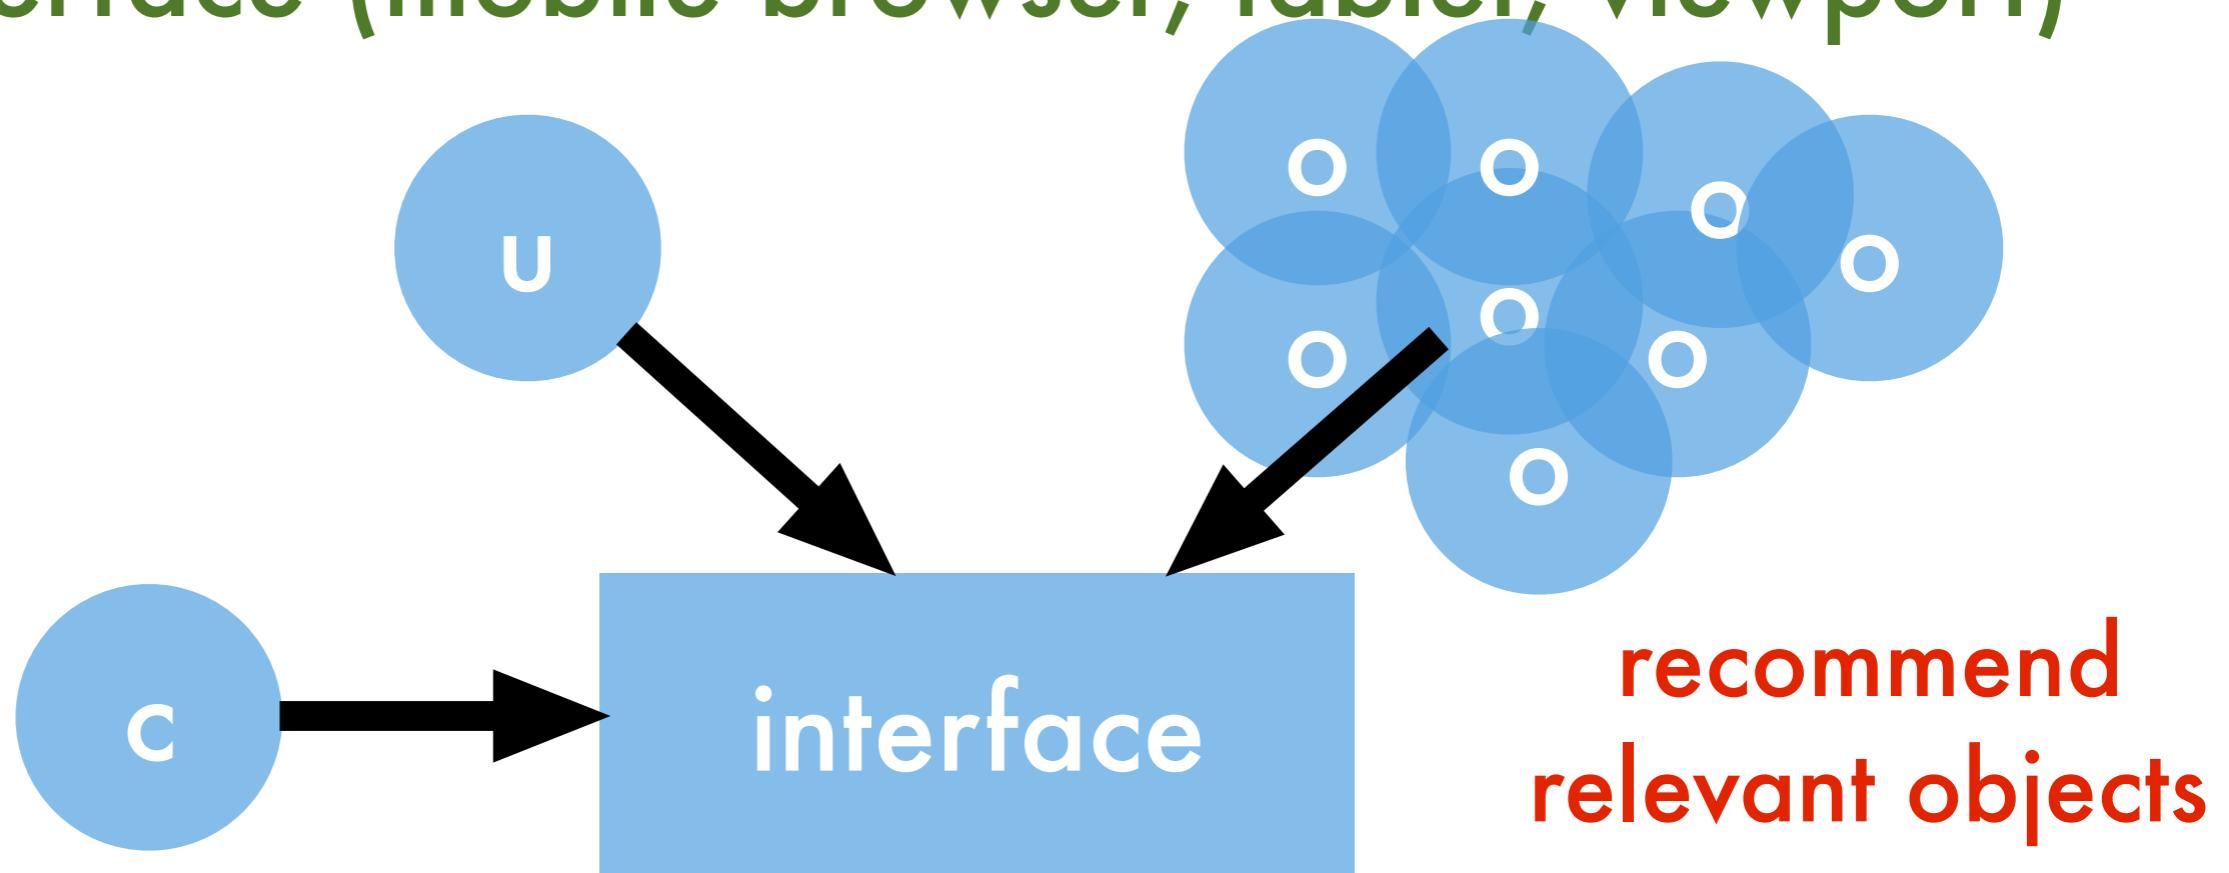
Average Customer Review: (1)

In Stock

Price: \$49.50[15 used & new from \\$45.49](#)[Add to Cart](#)[Add to Wish List](#) I own it Not interested Rate this itemRecommended because you purchased [Nonlinear Programming](#) ([Fix this](#))

A more formal view

- User (requests content)
- Objects (that can be displayed)
- Context (device, location, time)
- Interface (mobile browser, tablet, viewport)



Examples

- Movie recommendation (Netflix)
- Related product recommendation (Amazon)
- Web page ranking (Google)
- Social recommendation (Facebook)
- News content recommendation (Yahoo)
- Priority inbox & spam filtering (Google)
- Online dating (OK Cupid)
- Computational Advertising (Yahoo)

Running Example

Netflix Movie Recommendation

Training data

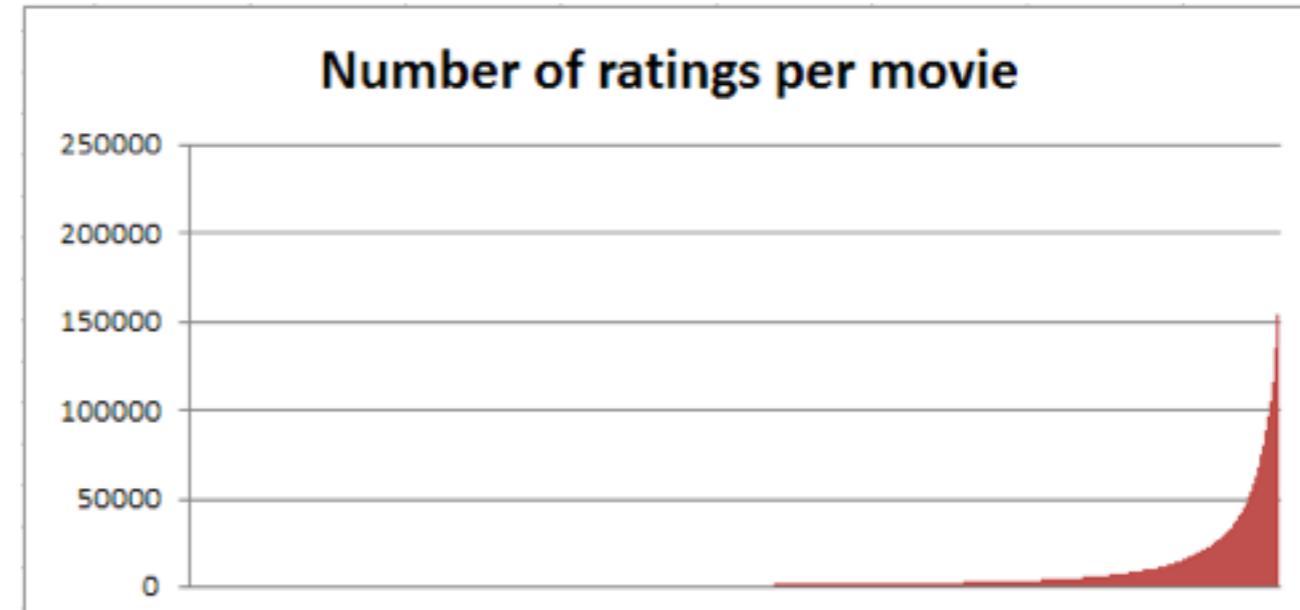
user	movie	date	score
1	21	5/7/02	1
1	213	8/2/04	5
2	345	3/6/01	4
2	123	5/1/05	4
2	768	7/15/02	3
3	76	1/22/01	5
4	45	8/3/00	4
5	568	9/10/05	1
5	342	3/5/03	2
5	234	12/28/00	2
6	76	8/11/02	5
6	56	6/15/03	4

Test data

user	movie	date	score
1	62	1/6/05	?
1	96	9/13/04	?
2	7	8/18/05	?
2	3	11/22/05	?
3	47	6/13/02	?
3	15	8/12/01	?
4	41	9/1/00	?
4	28	8/27/05	?
5	93	4/4/05	?
5	74	7/16/03	?
6	69	2/14/04	?
6	83	10/3/03	?

Challenges

- Scalability
 - Millions of objects
 - 100s of millions of users
 - Cold start
 - Changing user base
 - Changing inventory (movies, stories, goods)
 - Attributes
 - Imbalanced dataset
- User activity / item reviews
are power law distributed



Netflix competition yardstick

- Least mean squares prediction error
 - Easy to define

$$\text{rmse}(S) = \sqrt{|S|^{-1} \sum_{(i,u) \in S} (\hat{r}_{ui} - r_{ui})^2}$$

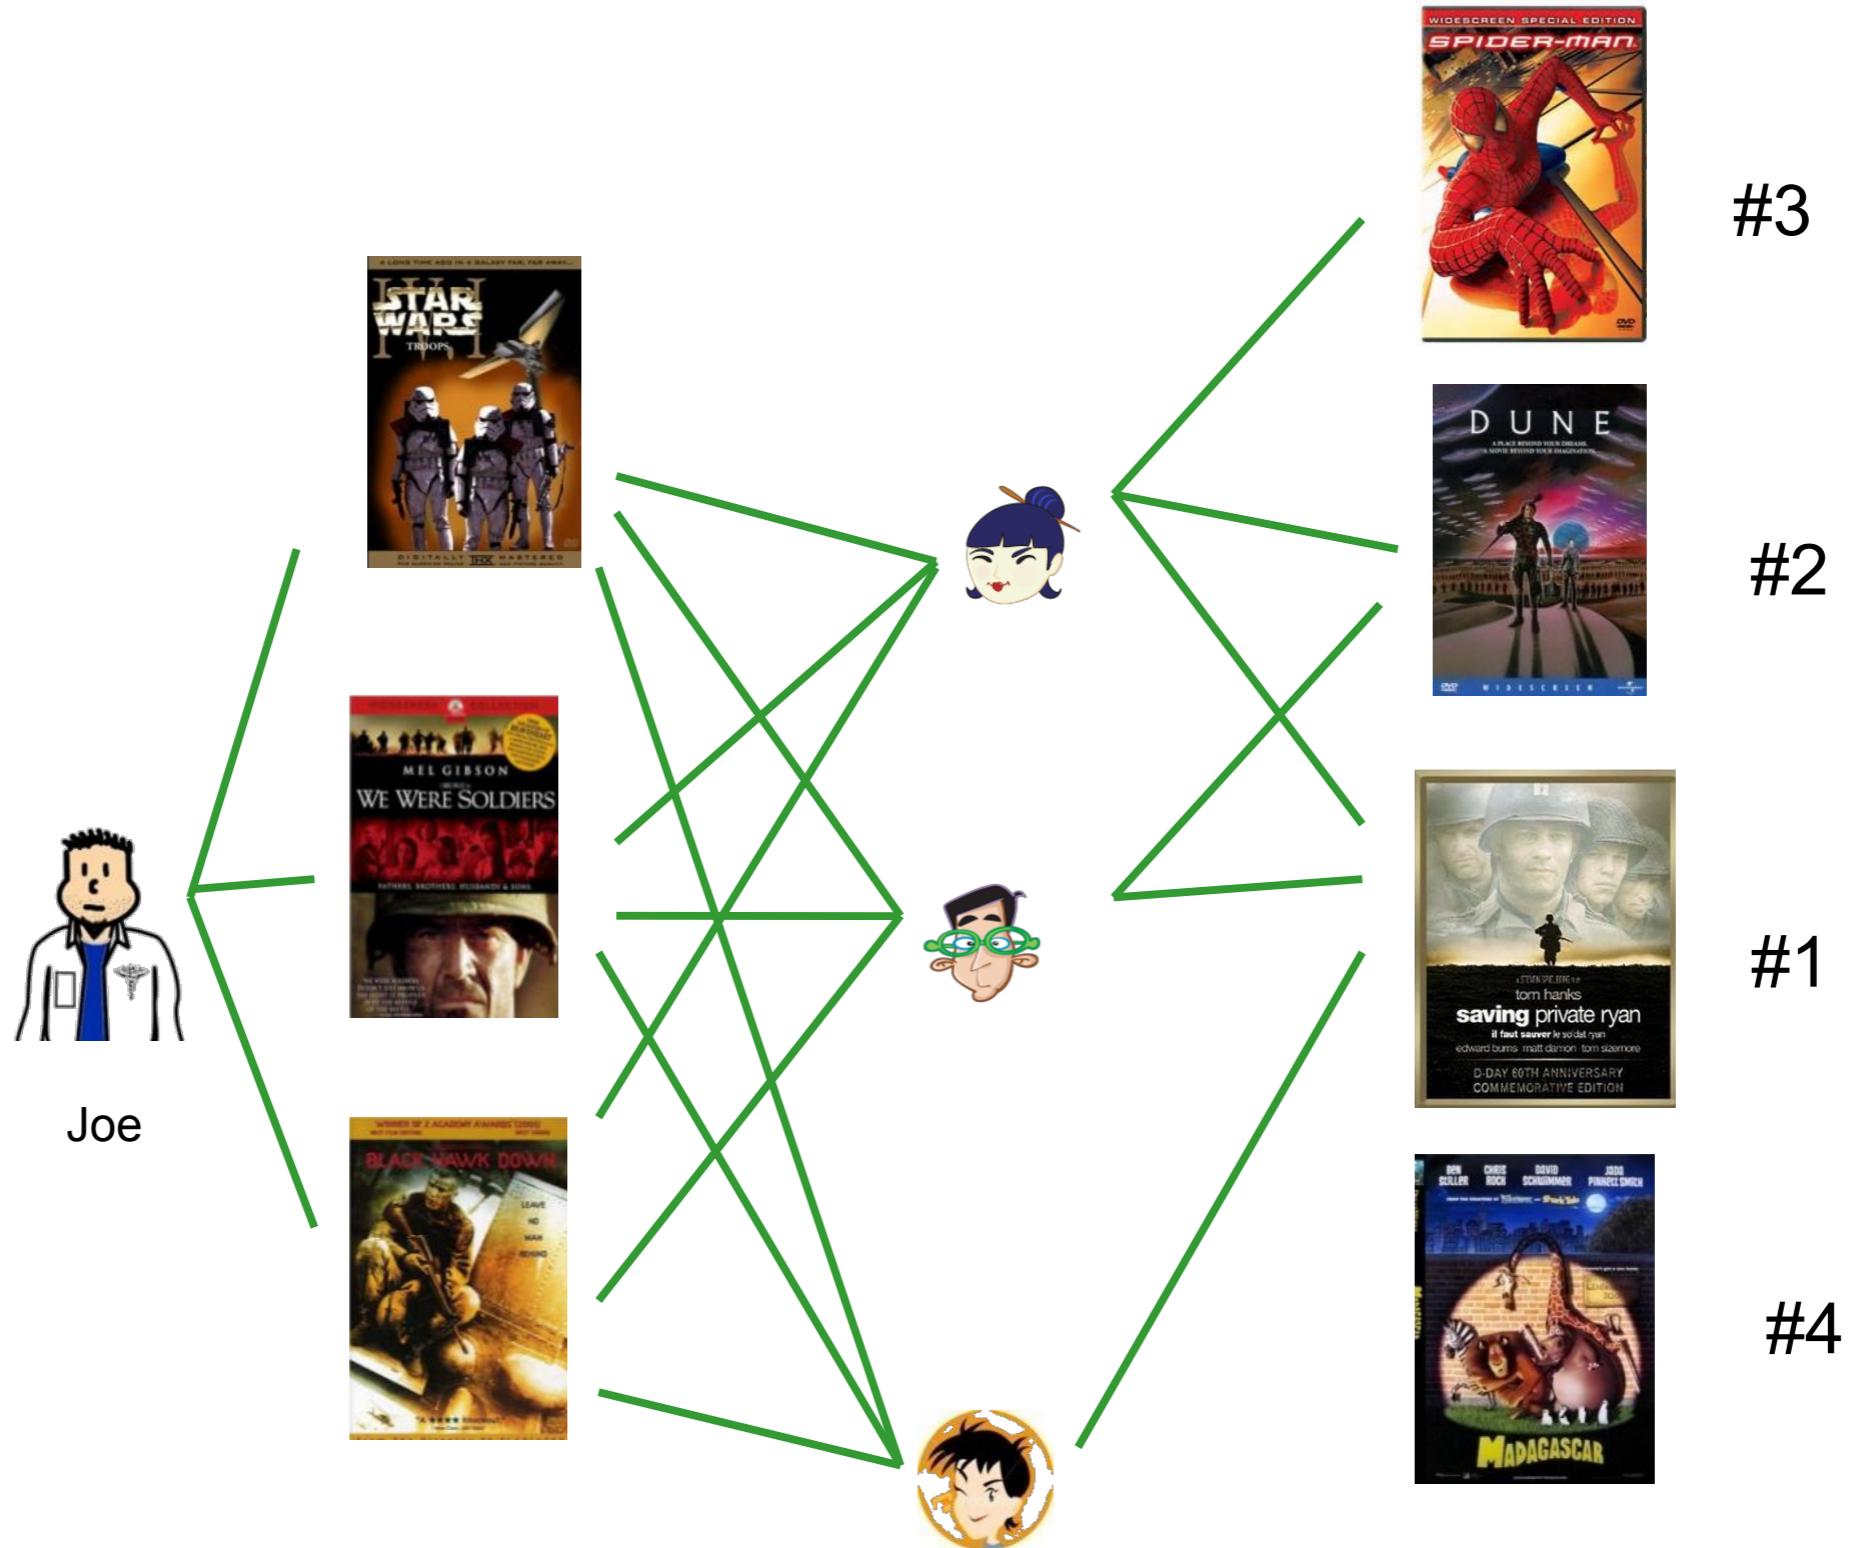
- Wrong measure for composing sessions!



- Consistent (in large sample size limit this will converge to minimizer)

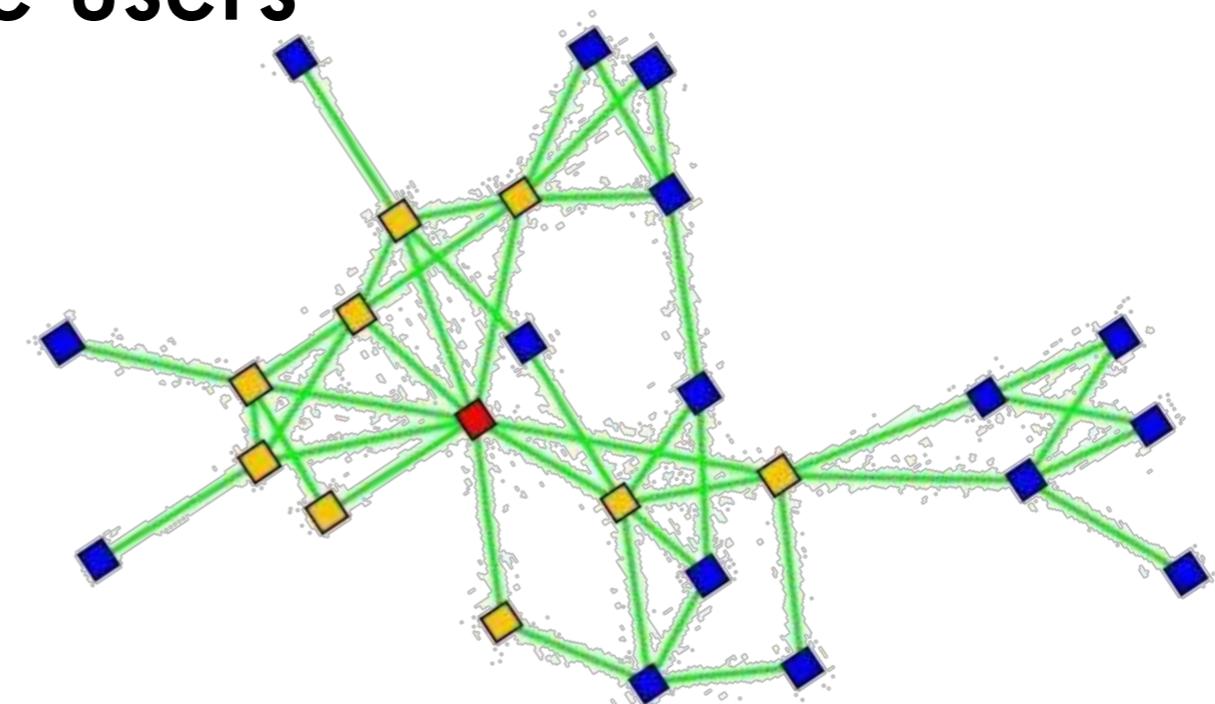
1 Neighborhood Methods

Basic Idea



Basic Idea

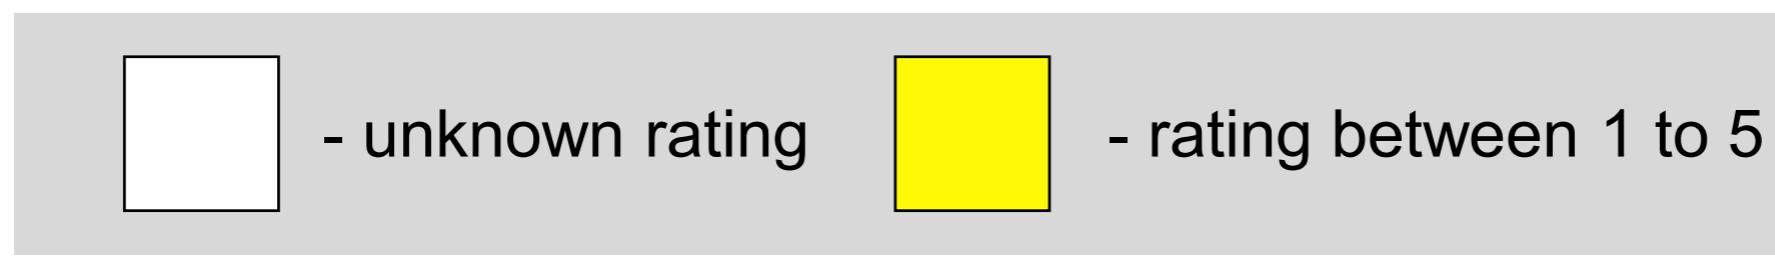
- (user,user) similarity to recommend items
 - good if item base is smaller than user base
 - good if item base changes rapidly
 - traverse bipartite similarity graph
- (item,item) similarity to recommend new items that were also liked by the same users
 - good if the user base is small
 - Oldest known CF method



Neighborhood based CF

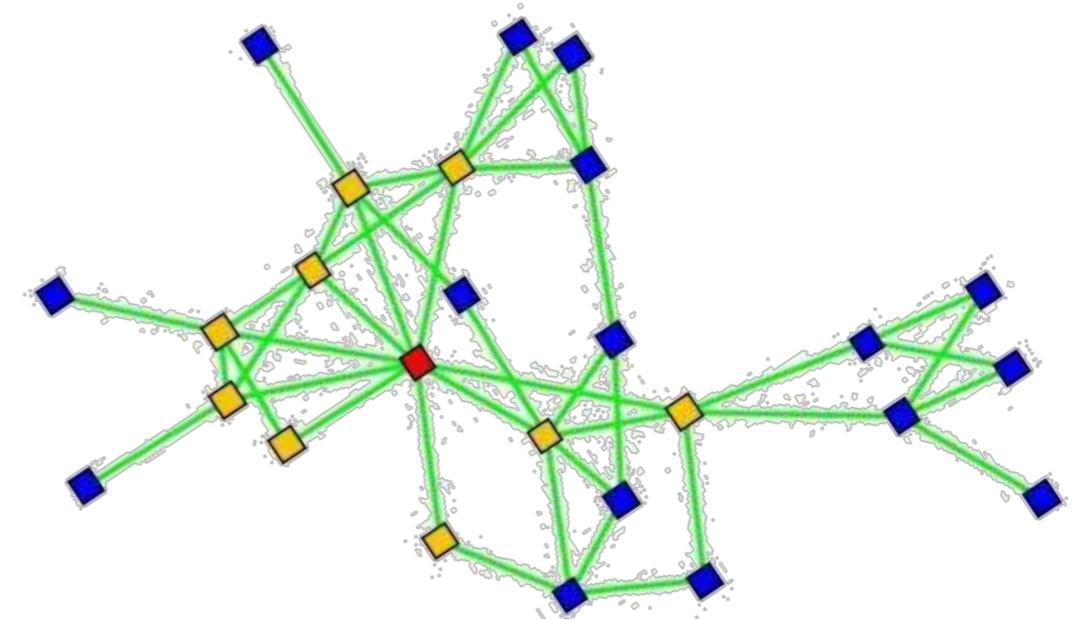
	users											
	1	2	3	4	5	6	7	8	9	10	11	12
1	1		3		2?6	5			5		4	
2				5	4			4		2	1	3
3	3	2	4		1	2		3		4	3	5
4		2	4		5			4			2	
5			4	3	4	2				2	5	
6	6	1		3		3		2			4	

similarity
 $s_{13} = 0.2$
 $s_{16} = 0.3$



Properties

- Intuitive
- No (substantial) training
- Handles new users / items
- Easy to explain to user



Recommended for you

+ Add as Playlist More

Casually Introducing Walter Smith III	Companeros De Mi Vida	Tibiri Tabara Sierra Maestra	New York Ska-Jazz Ensemble	More Late Night Transmissions With... Jaya the Cat	Appetite For Destruction
Similar to Eric Harland	Eliades Ochoa	You've scrobbled Sierra Maestra, but not this release	New York Ska-Jazz Ensemble	You've scrobbled Jaya the Cat, but not this release	Guns N' Roses
	Similar to Cachao and Irakere		You've scrobbled New York Ska-Jazz Ensemble,		You've scrobbled Guns N' Roses, but not this release

A blue arrow points from the bottom right towards the 'Appetite For Destruction' row, highlighting it.

Normalization / Bias

- Problem
 - Some items are significantly higher rated
 - Some users rate substantially lower
 - Ratings change over time
- Bias correction is crucial for nearest neighborhood recommender algorithm
 - Offset per user
 - Offset per movie
 - Time effects
 - Global bias

$$b_{ui} = \mu + b_u + b_i$$

Baseline estimation

- Mean rating is 3.7
- Troll Hunter is 0.7 above mean
- User rates 0.2 below mean
- Baseline is 4.2 stars
- Least mean squares problem

$$\underset{b}{\text{minimize}} \sum_{(u,i)} (r_{ui} - \mu - b_u - b_i)^2 + \lambda \left[\sum_u b_u^2 + \sum_i b_i^2 \right]$$



- Jointly convex. Alternatively remove mean & iterate

$$b_i = \frac{\sum_{u \in R(i)} (r_{ui} - \mu - b_u)}{\lambda + |R(i)|} \quad \text{and} \quad b_u = \frac{\sum_{i \in R(u)} (r_{ui} - \mu - b_i)}{\lambda + |R(u)|}$$

Parzen Windows style CF

- Similarity measure s_{ij} between items
- Find set $s_k(i,u)$ of k-nearest neighbors to i that were rated by user u
- Weighted average over the set

$$\hat{r}_{ui} = b_{ui} + \frac{\sum_{j \in s_k(i,u)} s_{ij} (r_{uj} - b_{uj})}{\sum_{j \in s_k(i,u)} s_{ij}} \text{ where } b_{ui} = \mu + b_u + b_i$$

- How to compute s_{ij} ?

(item,item) similarity measures

User ratings for item **i**:

1	?	?	?	5	5	3	?	?	?	4	2	?	?	?	?	4	?	5	4	1	?
---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---

User ratings for item **j**:

?	?	4	2	5	?	?	1	2	5	?	?	2	?	?	3	?	?	?	5	4
---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---

- Pearson correlation coefficient
 - nonuniform support
 - compute only over shared support
 - shrinkage towards 0 to address problem of small support (typically few items in

(item,item) similarity

- Empirical Pearson correlation coefficient

$$\hat{\rho}_{ij} = \frac{\sum_{u \in U(i,j)} (r_{ui} - b_{ui})(r_{uj} - b_{uj})}{\sqrt{\sum_{u \in U(i,j)} (r_{ui} - b_{ui})^2 \sum_{u \in U(i,j)} (r_{uj} - b_{uj})^2}}$$

- Smoothing towards 0 for small support

$$s_{ij} = \frac{|U(i,j)| - 1}{|U(i,j)| - 1 + \lambda} \hat{\rho}_{ij}$$

- Here, $|U(i,j)|$ users that rated both i and j
- Shrink towards baseline for small neighborhood

$$\hat{r}_{ui} = b_{ui} + \frac{\sum_{j \in s_k(i,u)} s_{ij} (r_{uj} - b_{uj})}{\lambda + \sum_{j \in s_k(i,u)} s_{ij}}$$

Similarity for binary data

- Pearson correlation meaningless

- Views

m_i users acting on i

- Purchase behavior

m_{ij} users acting on both i and j

- Clicks

m total number of users

- Jaccard similarity

(intersection vs. joint) $s_{ij} = \frac{m_{ij}}{\alpha + m_i + m_j - m_{ij}}$

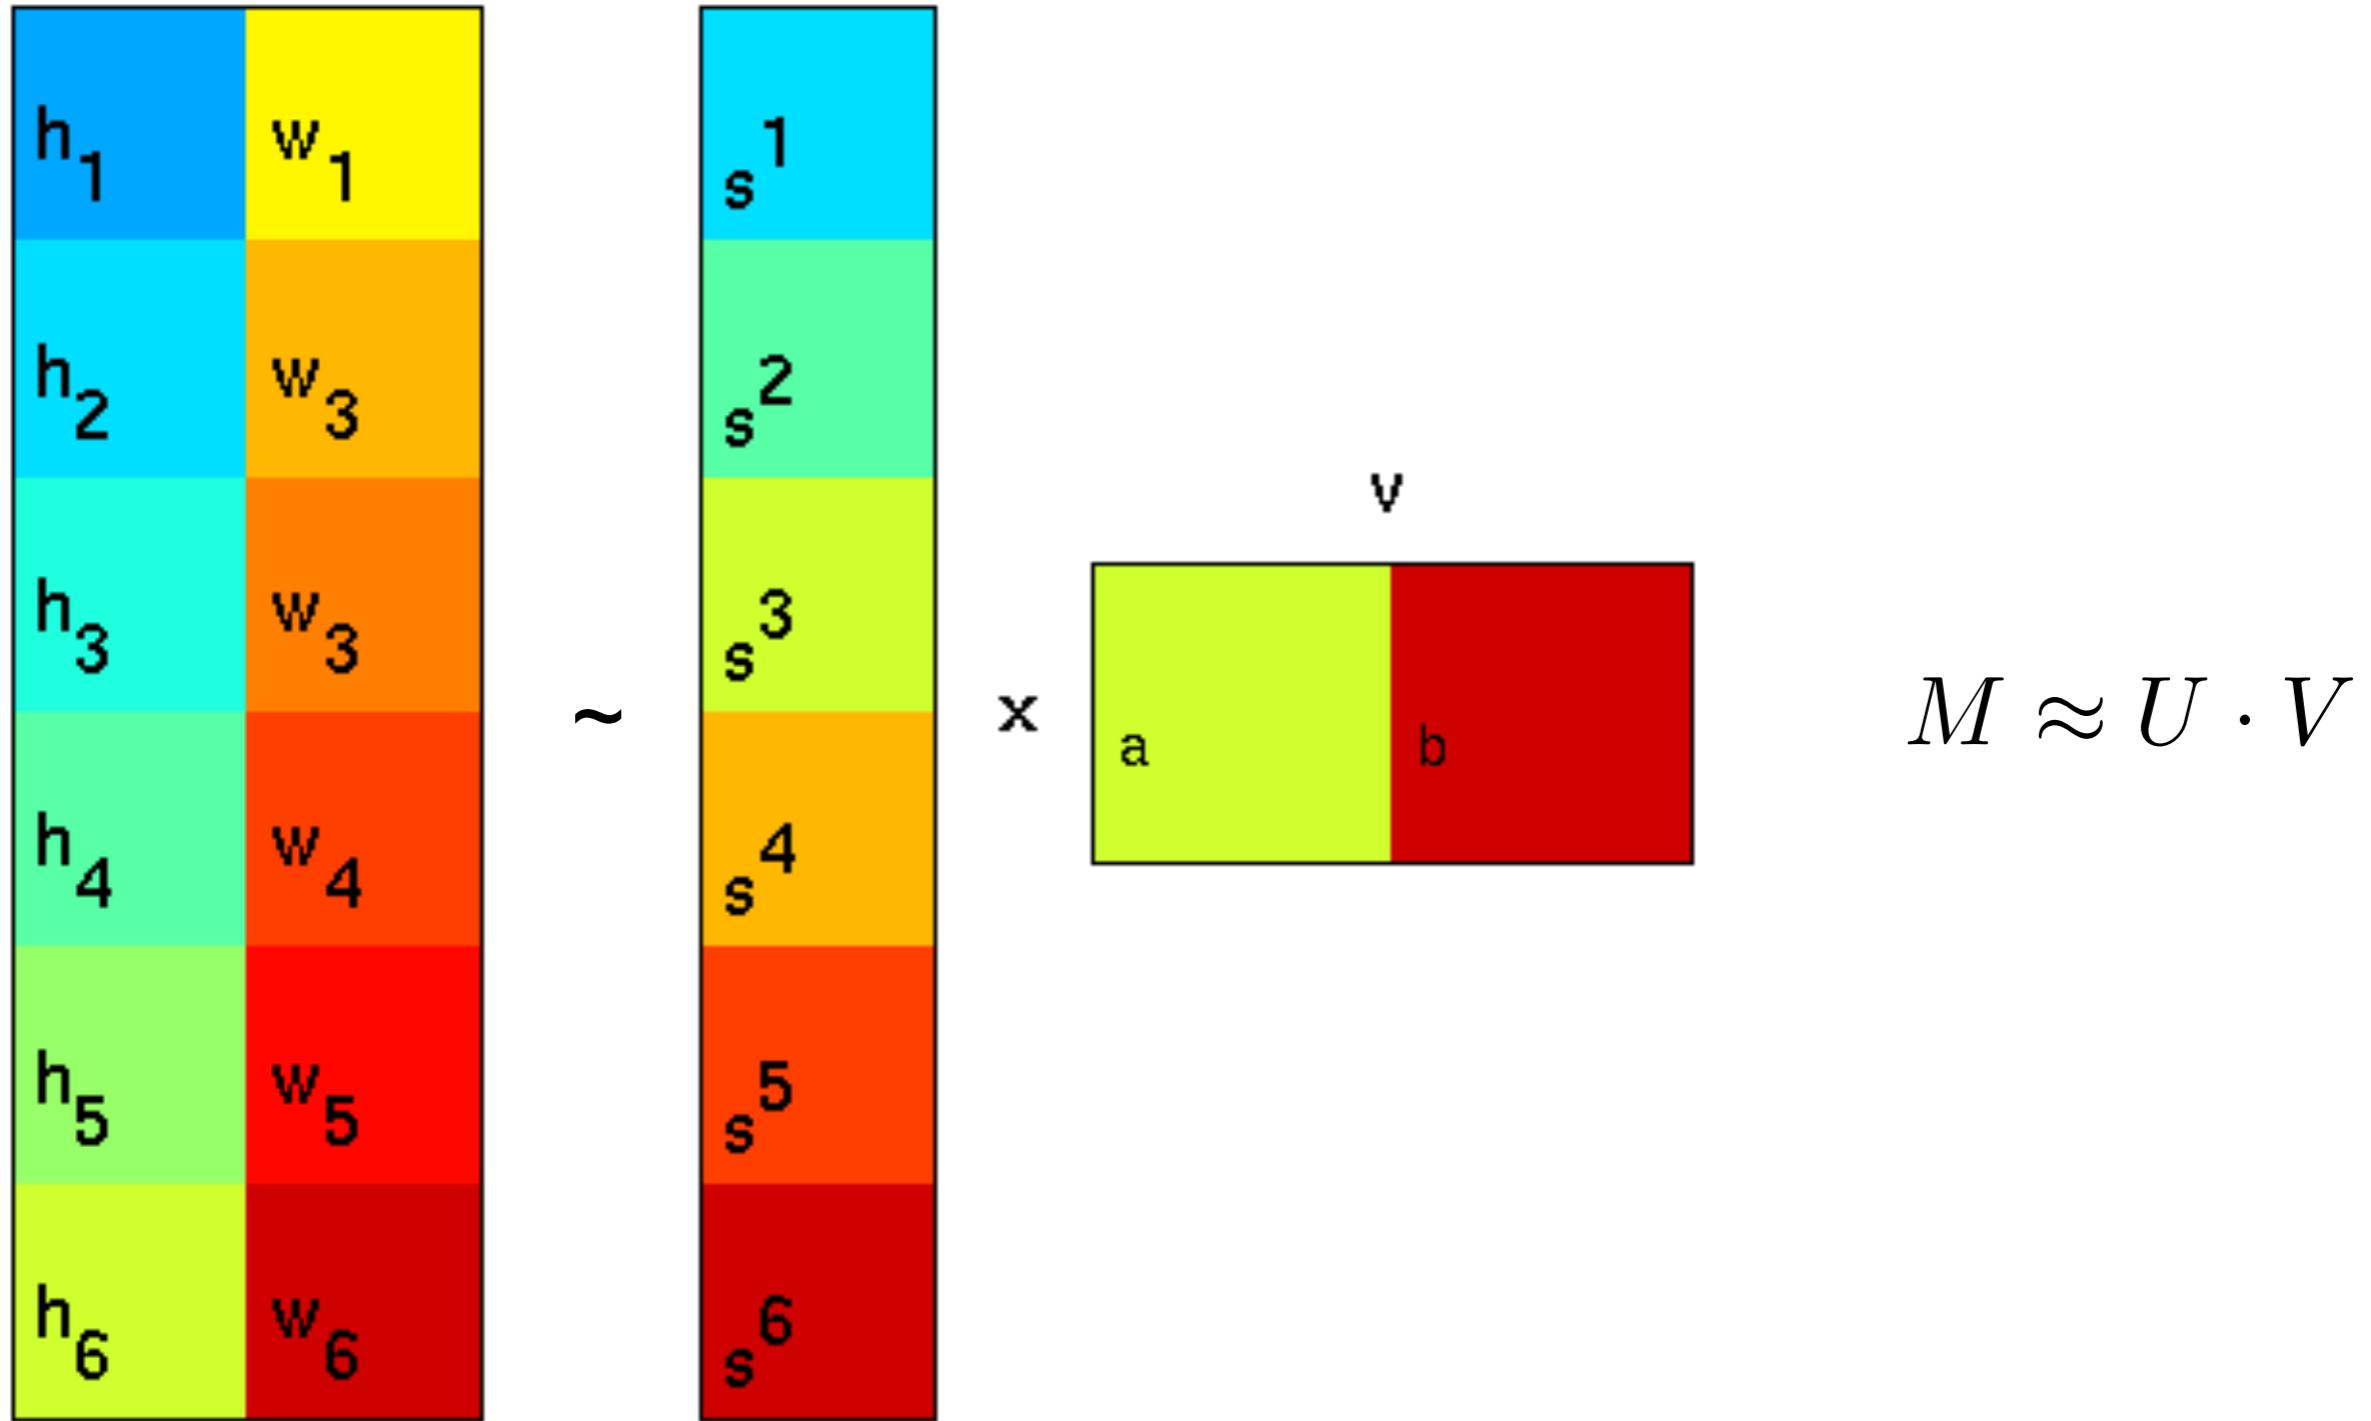
- Observed/expected ratio

Improve by counting $s_{ij} = \frac{\text{observed}}{\text{expected}} \approx \frac{m_{ij}}{\alpha + m_i m_j / m}$
per user (many users better than heavy users)

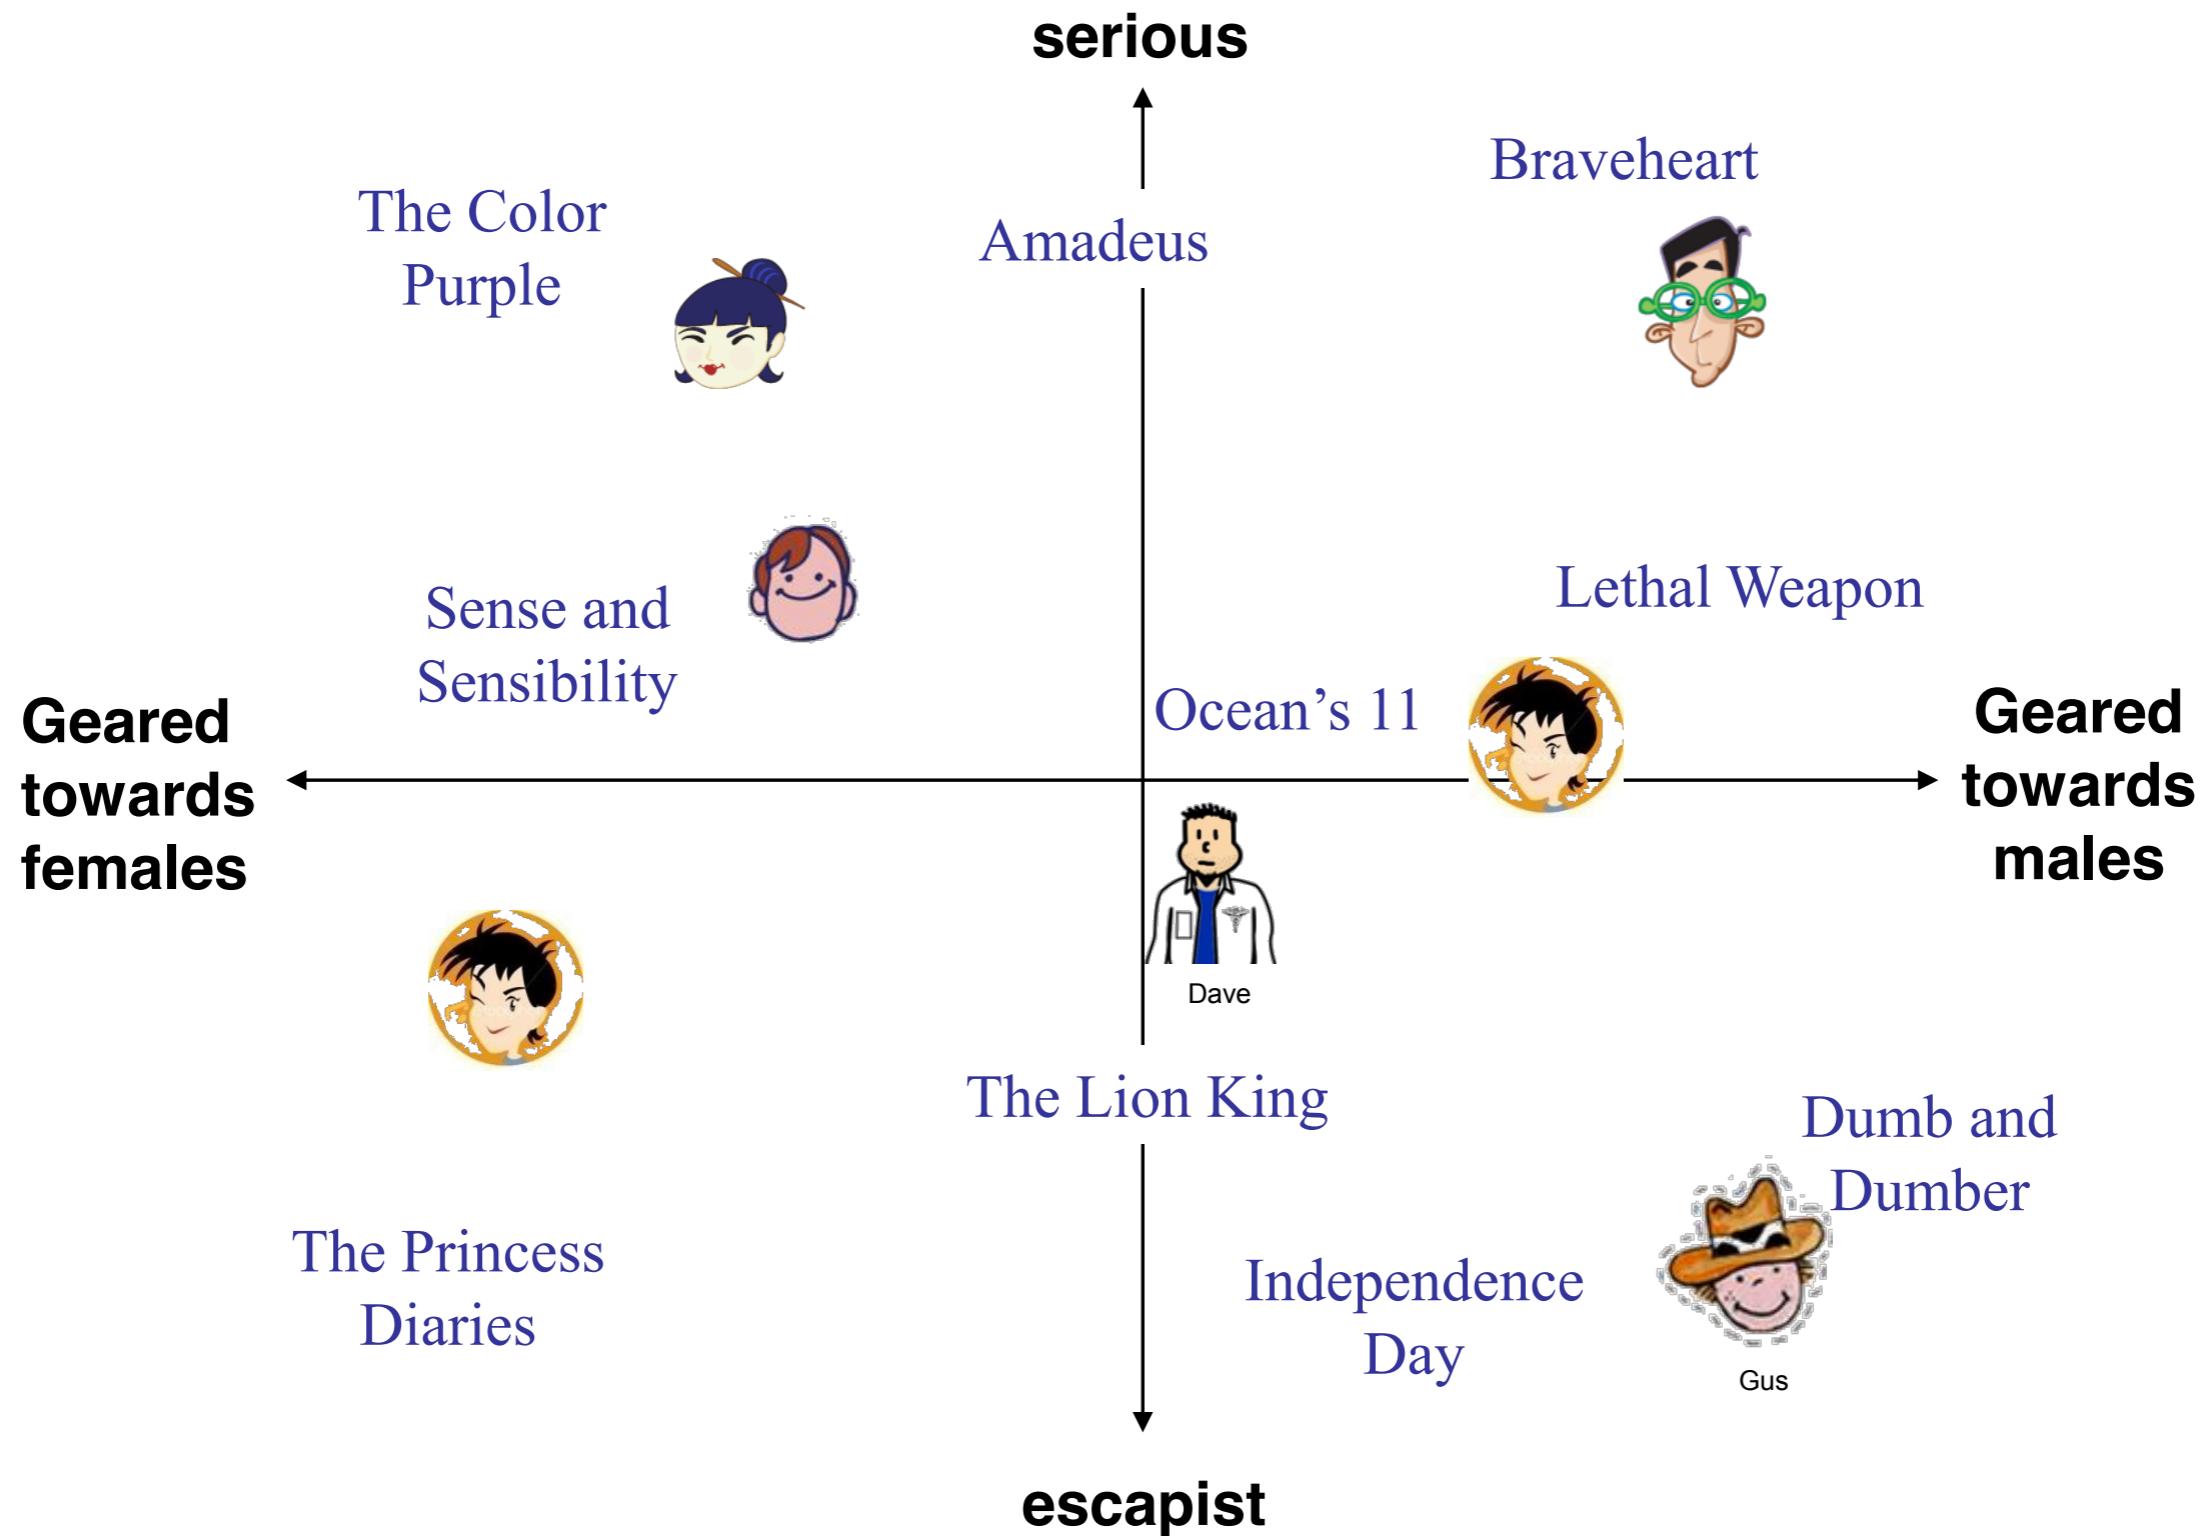
2 Matrix Factorization

Basics

Basic Idea



Latent variable view



Basic matrix factorization

users									
items	1		3		5		5	4	
			5	4		4		2	1
	2	4		1	2		3	4	3
		2	4		5		4		2
			4	3	4	2			2
	1		3		3		2		4

~

items	.1	-.4	.2
	-.5	.6	.5
	-.2	.3	.5
	1.1	2.1	.3
	-.7	2.1	-2
	-1	.7	.3

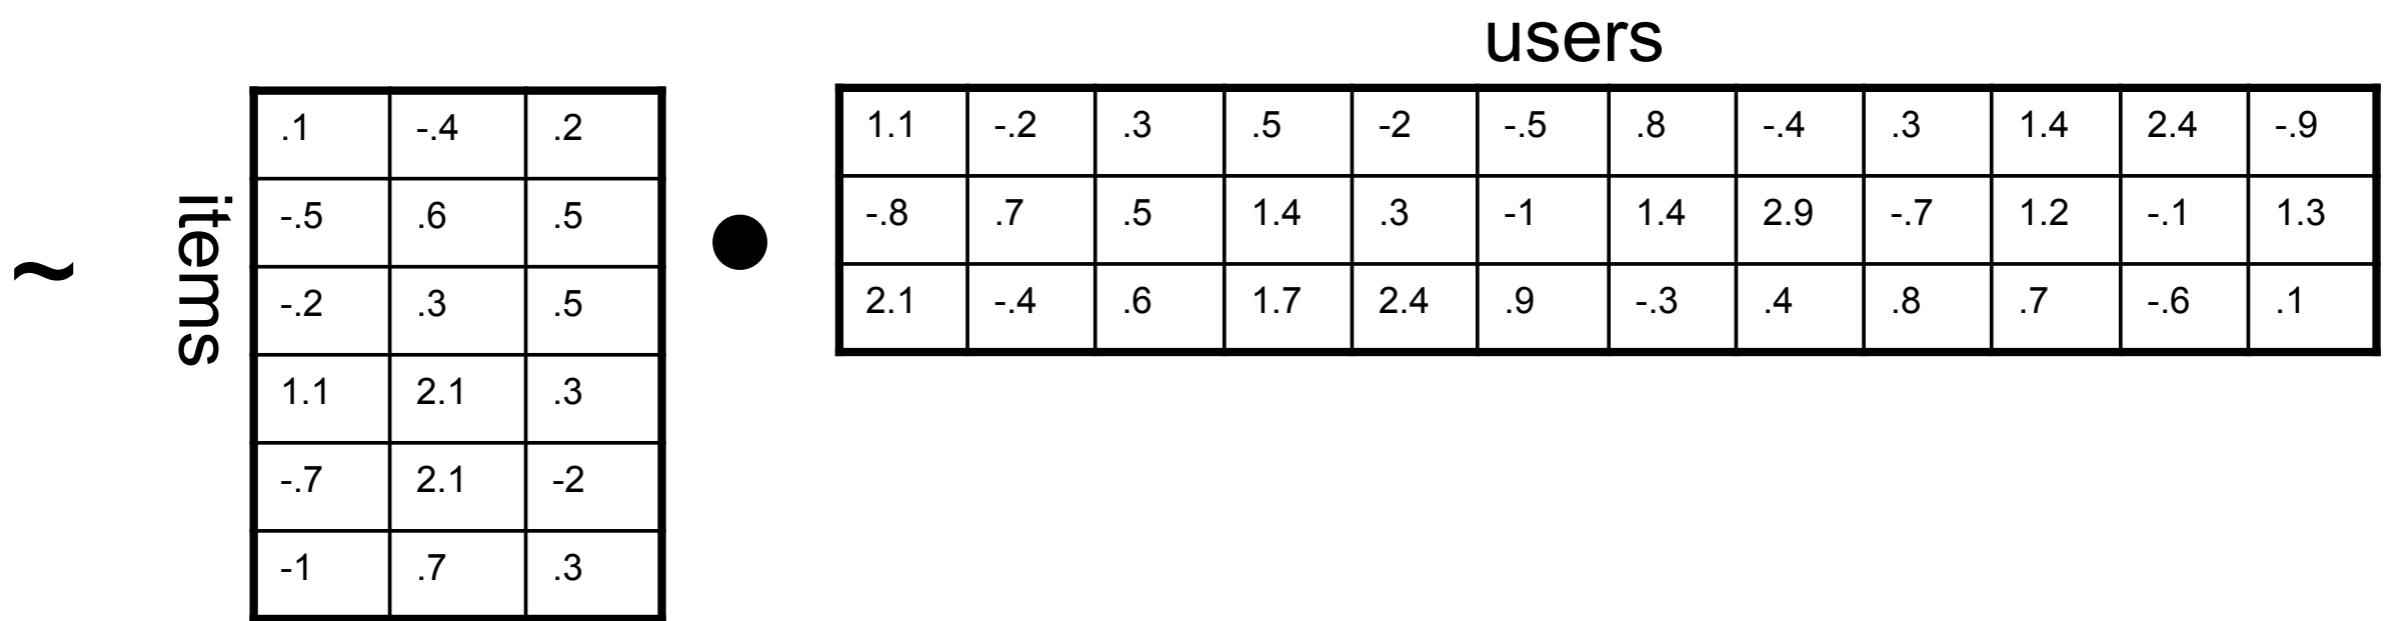
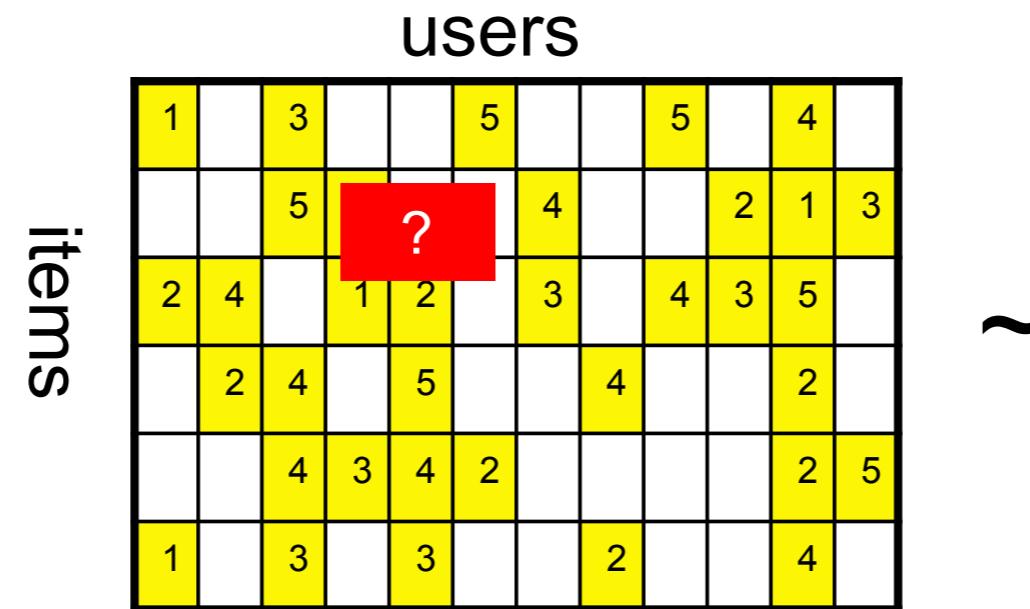
users

.1	-.2	.3	.5	-2	-.5	.8	-.4	.3	1.4	2.4	-.9
-.8	.7	.5	1.4	.3	-1	1.4	2.9	-.7	1.2	-.1	1.3
2.1	-.4	.6	1.7	2.4	.9	-.3	.4	.8	.7	-.6	.1

~

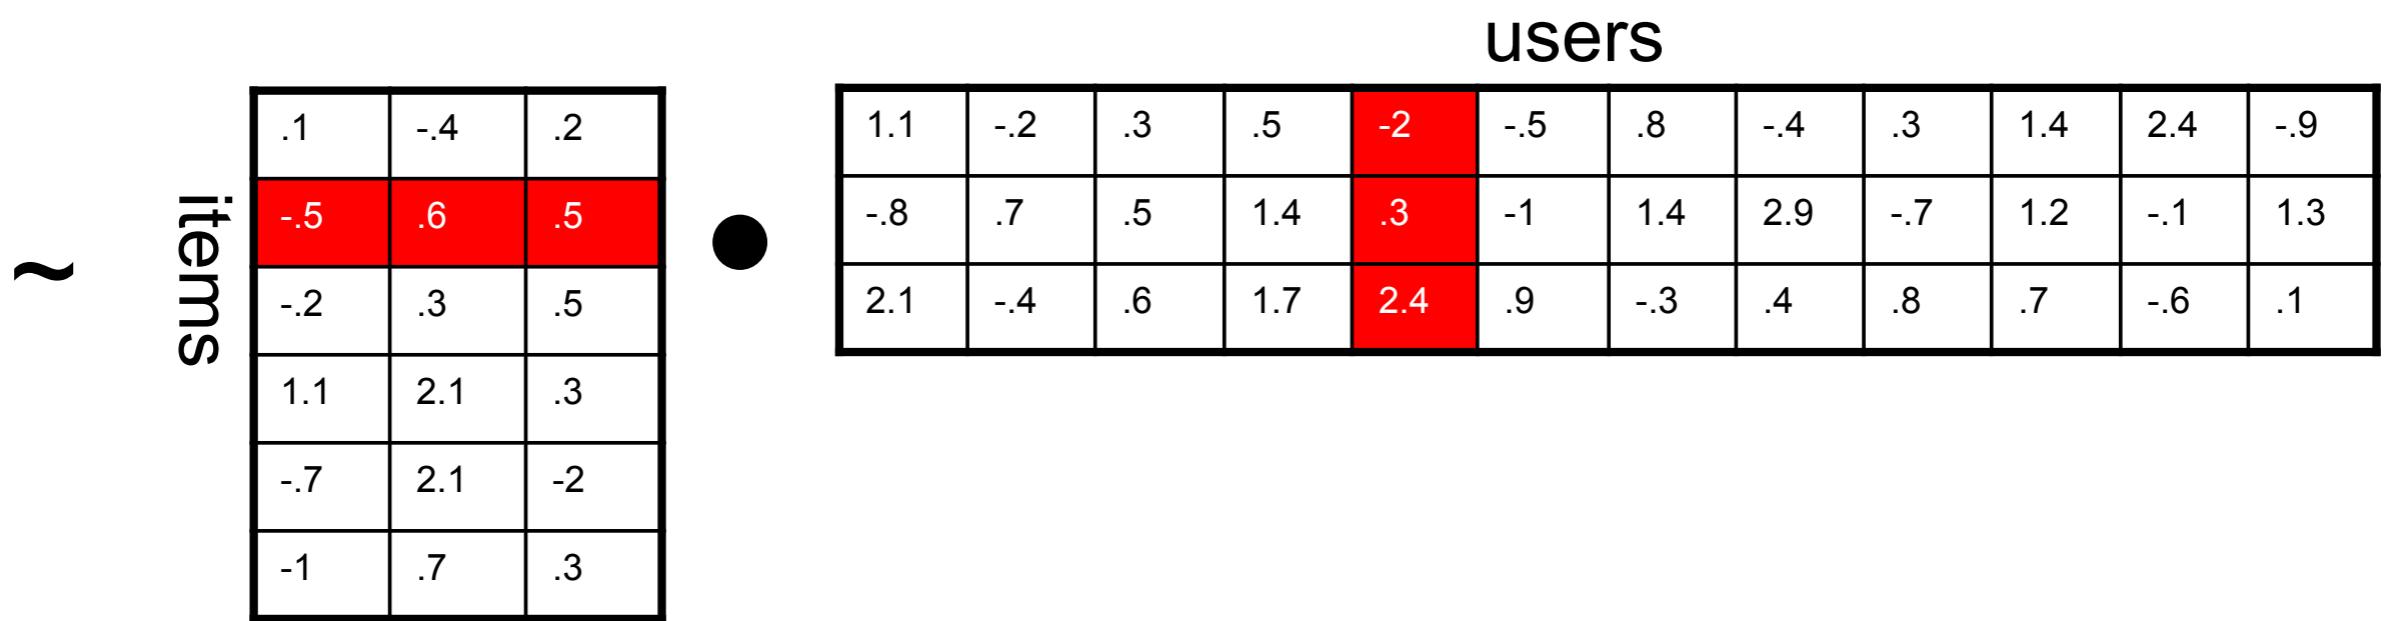
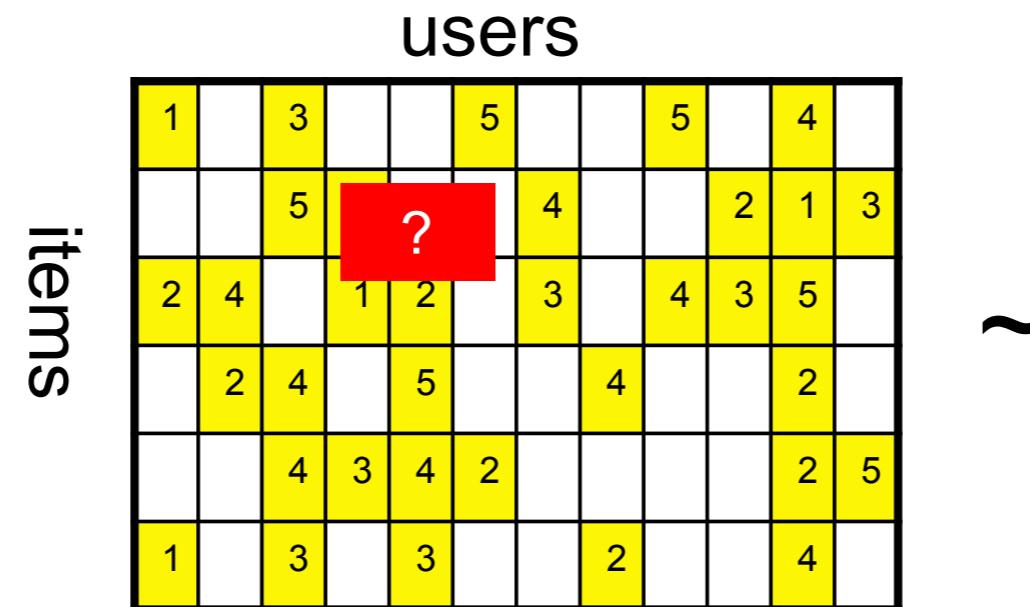
A rank-3 SVD approximation

Estimate unknown ratings as inner products of latent factors



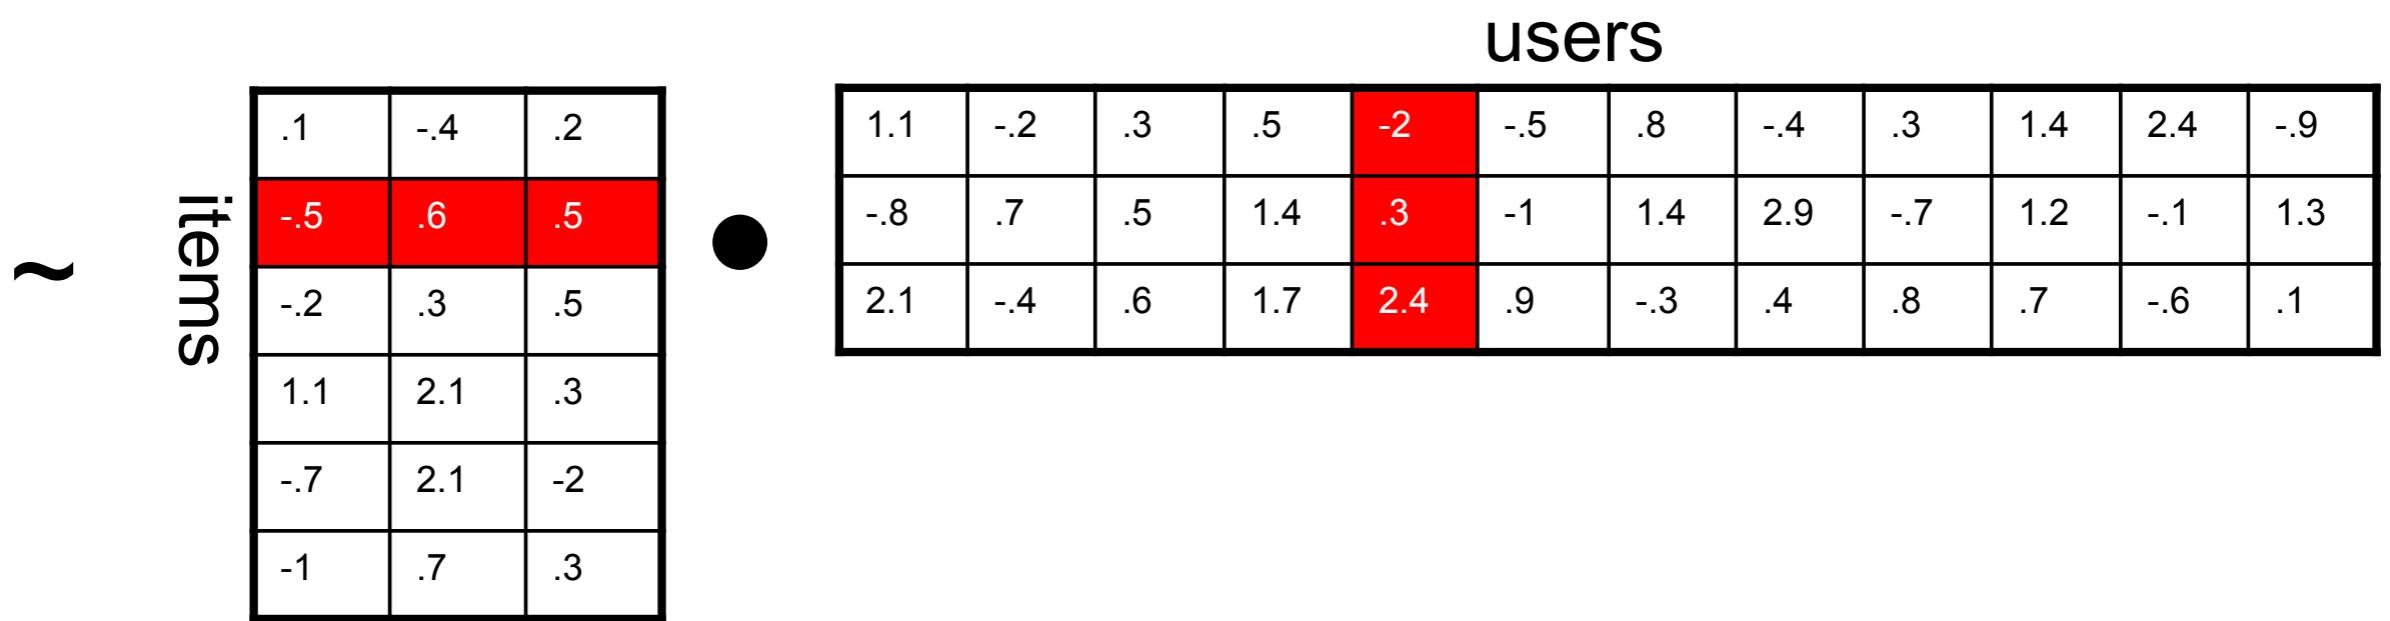
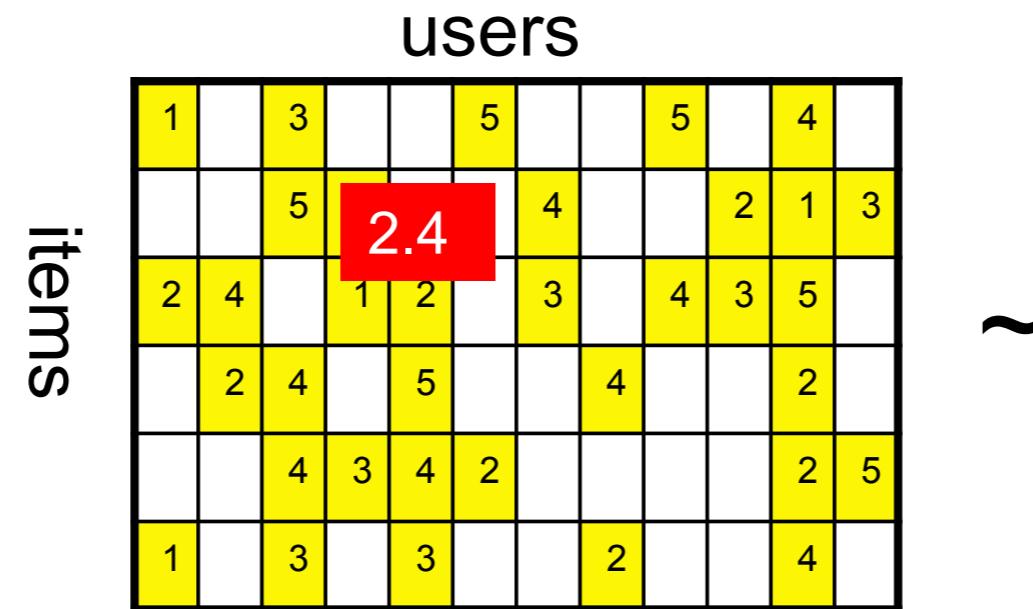
A rank-3 SVD approximation

Estimate unknown ratings as inner products of latent factors



A rank-3 SVD approximation

Estimate unknown ratings as inner products of latent factors



A rank-3 SVD approximation

Properties

1		3			5		5		4	
		5	4			4		2	1	3
2	4		1	2		3	4	3	5	
	2	4		5		4			2	
		4	3	4	2				2	5
1		3	3			2		4		

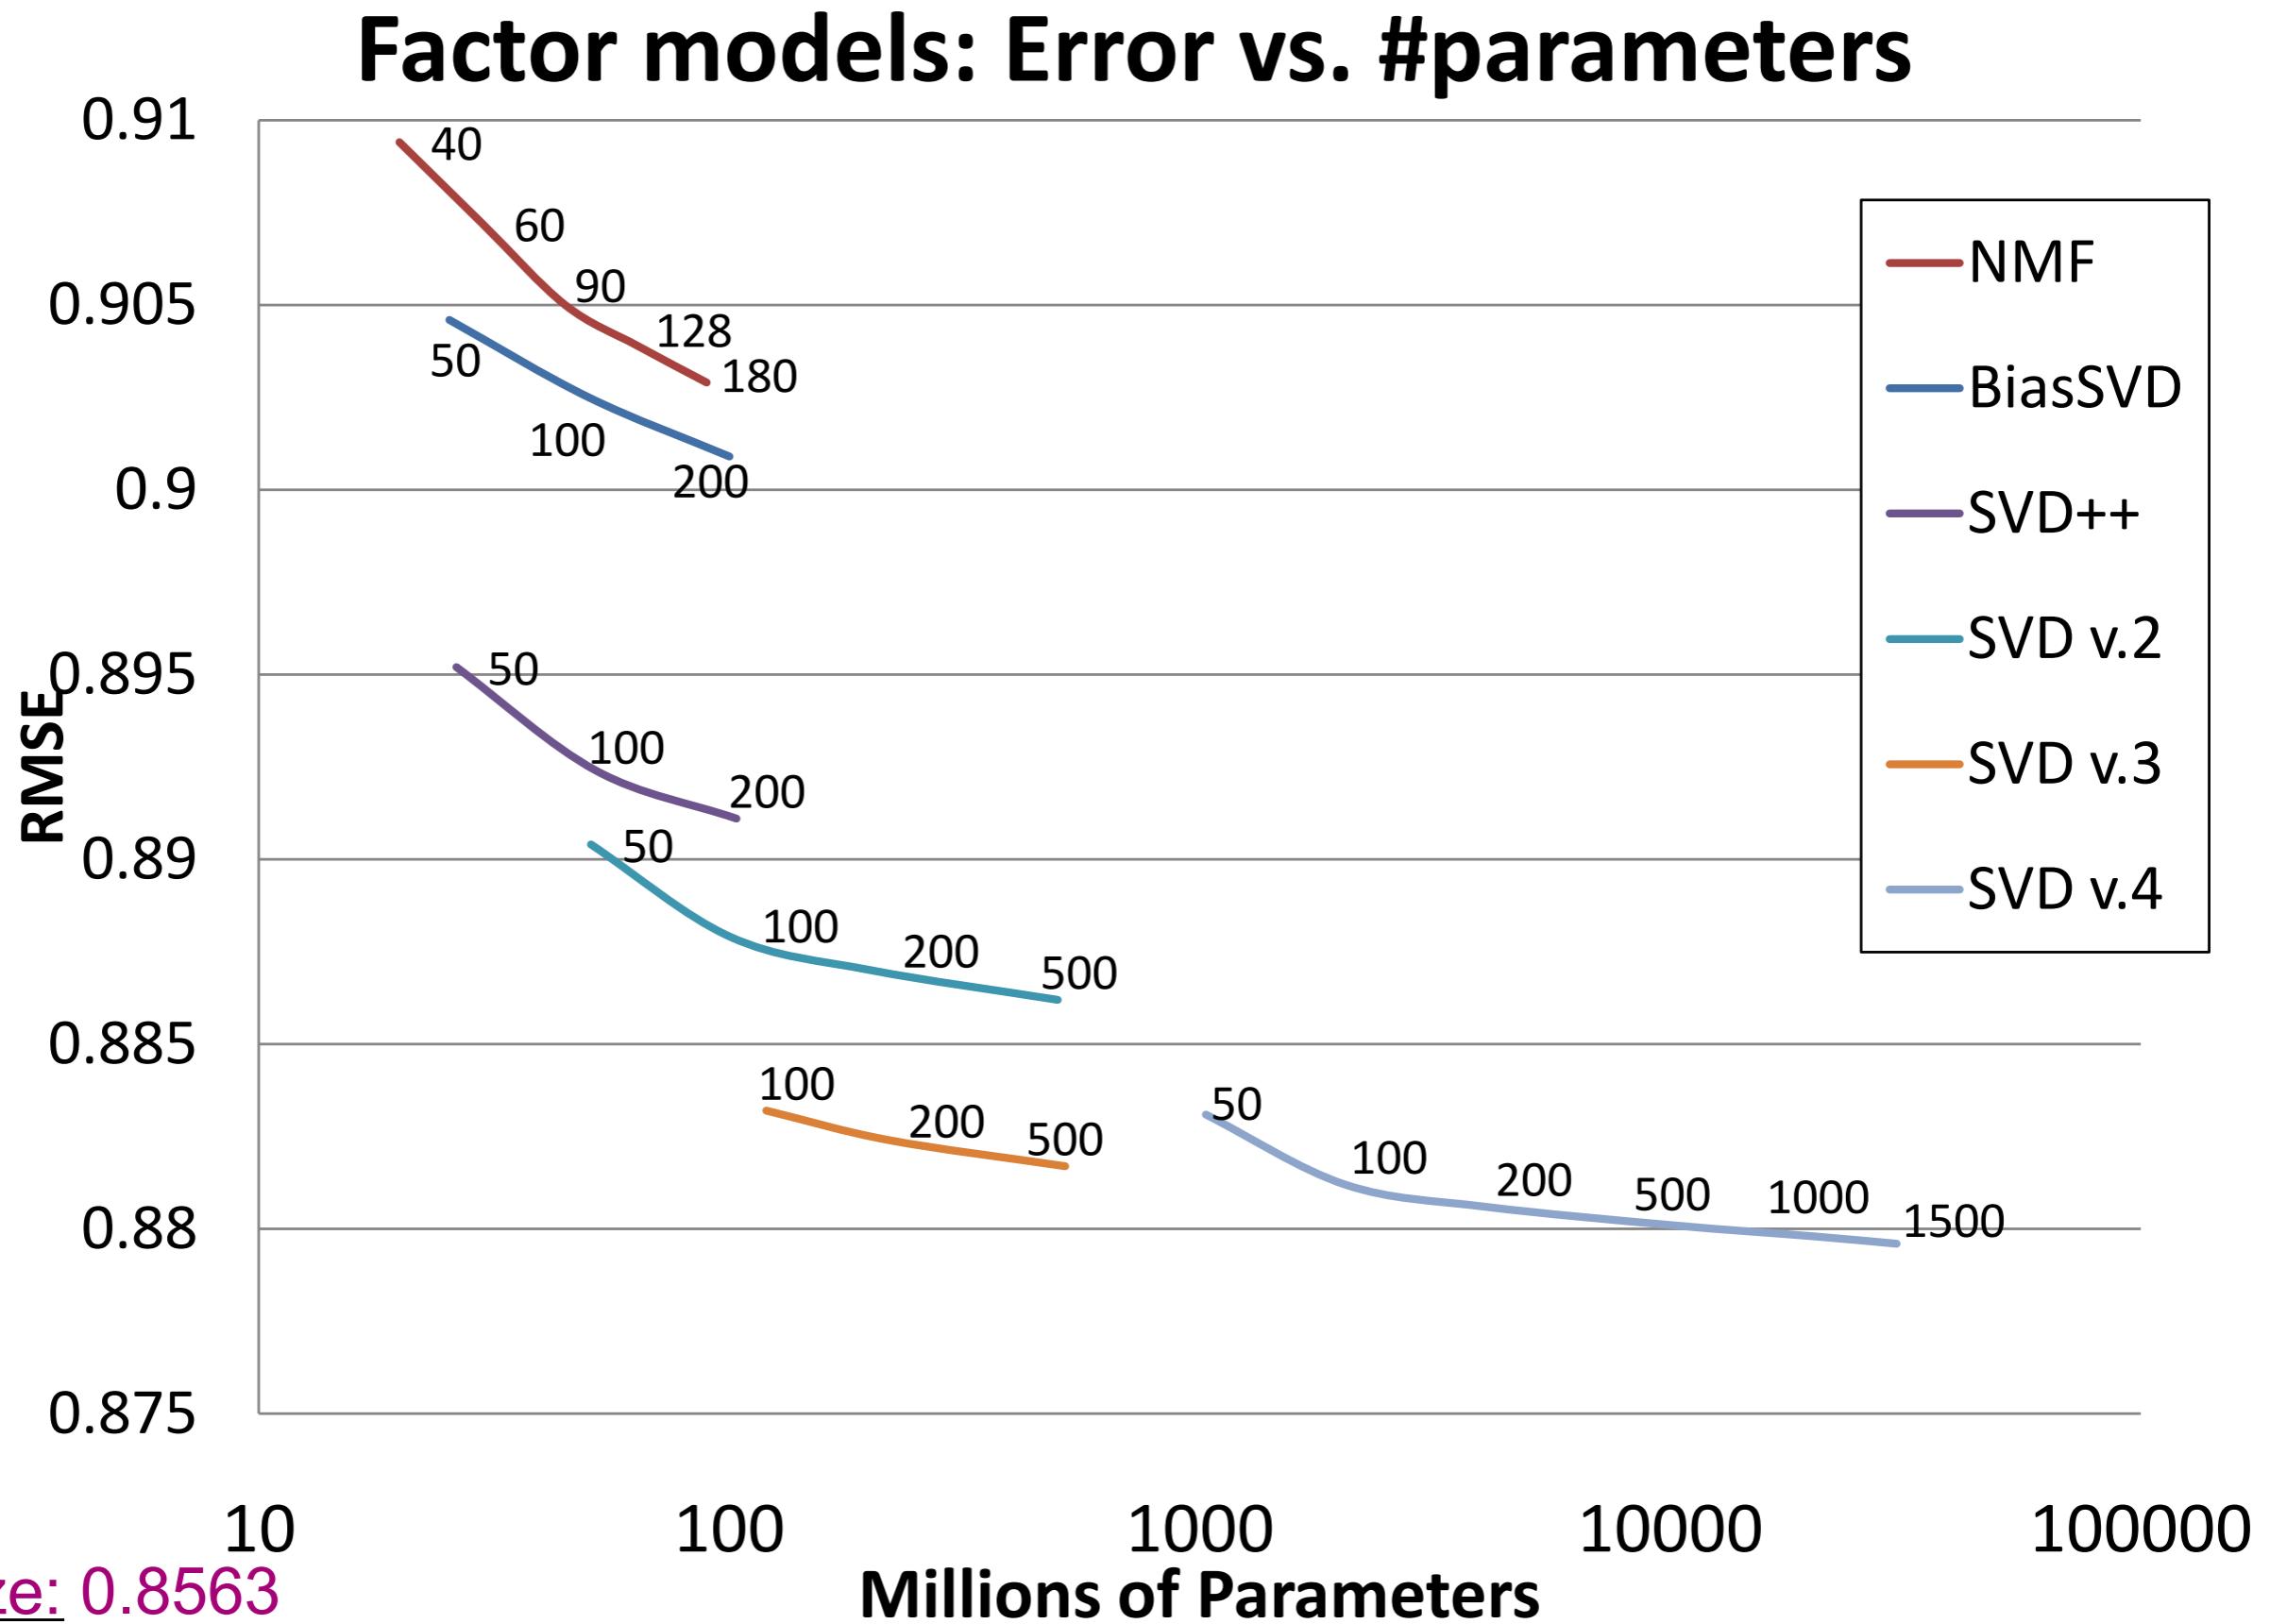
~

.1	-.4	.2
-.5	.6	.5
-.2	.3	.5
1.1	2.1	.3
-.7	2.1	-2
-1	.7	.3

1.1	-.2	.3	.5	-2	-.5	.8	-.4	.3	1.4	2.4	-.9
-.8	.7	.5	1.4	.3	-1	1.4	2.9	-.7	1.2	-.1	1.3
2.1	-.4	.6	1.7	2.4	.9	-.3	.4	.8	.7	-.6	.1

- SVD is undefined for missing entries
 - stochastic gradient descent (faster)
 - alternating optimization
- Overfitting without regularization particularly if fewer reviews than dimensions
- Very popular on Netflix

Netflix: 0.9514



Prize: 0.8563

Risk Minimization View

- **Objective Function**

$$\underset{p,q}{\text{minimize}} \sum_{(u,i) \in S} (r_{ui} - \langle p_u, q_i \rangle)^2 + \lambda \left[\|p\|_{\text{Frob}}^2 + \|q\|_{\text{Frob}}^2 \right]$$

- **Alternating least squares**

$$p_u \leftarrow \left[\lambda \mathbf{1} + \sum_{i | (u,i) \in S} q_i q_i^\top \right]^{-1} \sum_i q_i r_{ui}$$

good for
MapReduce

$$q_i \leftarrow \left[\lambda \mathbf{1} + \sum_{u | (u,i) \in S} p_u p_u^\top \right]^{-1} \sum_i p_u r_{ui}$$

Risk Minimization View

- **Objective Function**

$$\underset{p,q}{\text{minimize}} \sum_{(u,i) \in S} (r_{ui} - \langle p_u, q_i \rangle)^2 + \lambda \left[\|p\|_{\text{Frob}}^2 + \|q\|_{\text{Frob}}^2 \right]$$

- **Stochastic gradient descent**

$$p_u \leftarrow (1 - \lambda \eta_t) p_u - \eta_t q_i (r_{ui} - \langle p_u, q_i \rangle)$$

much

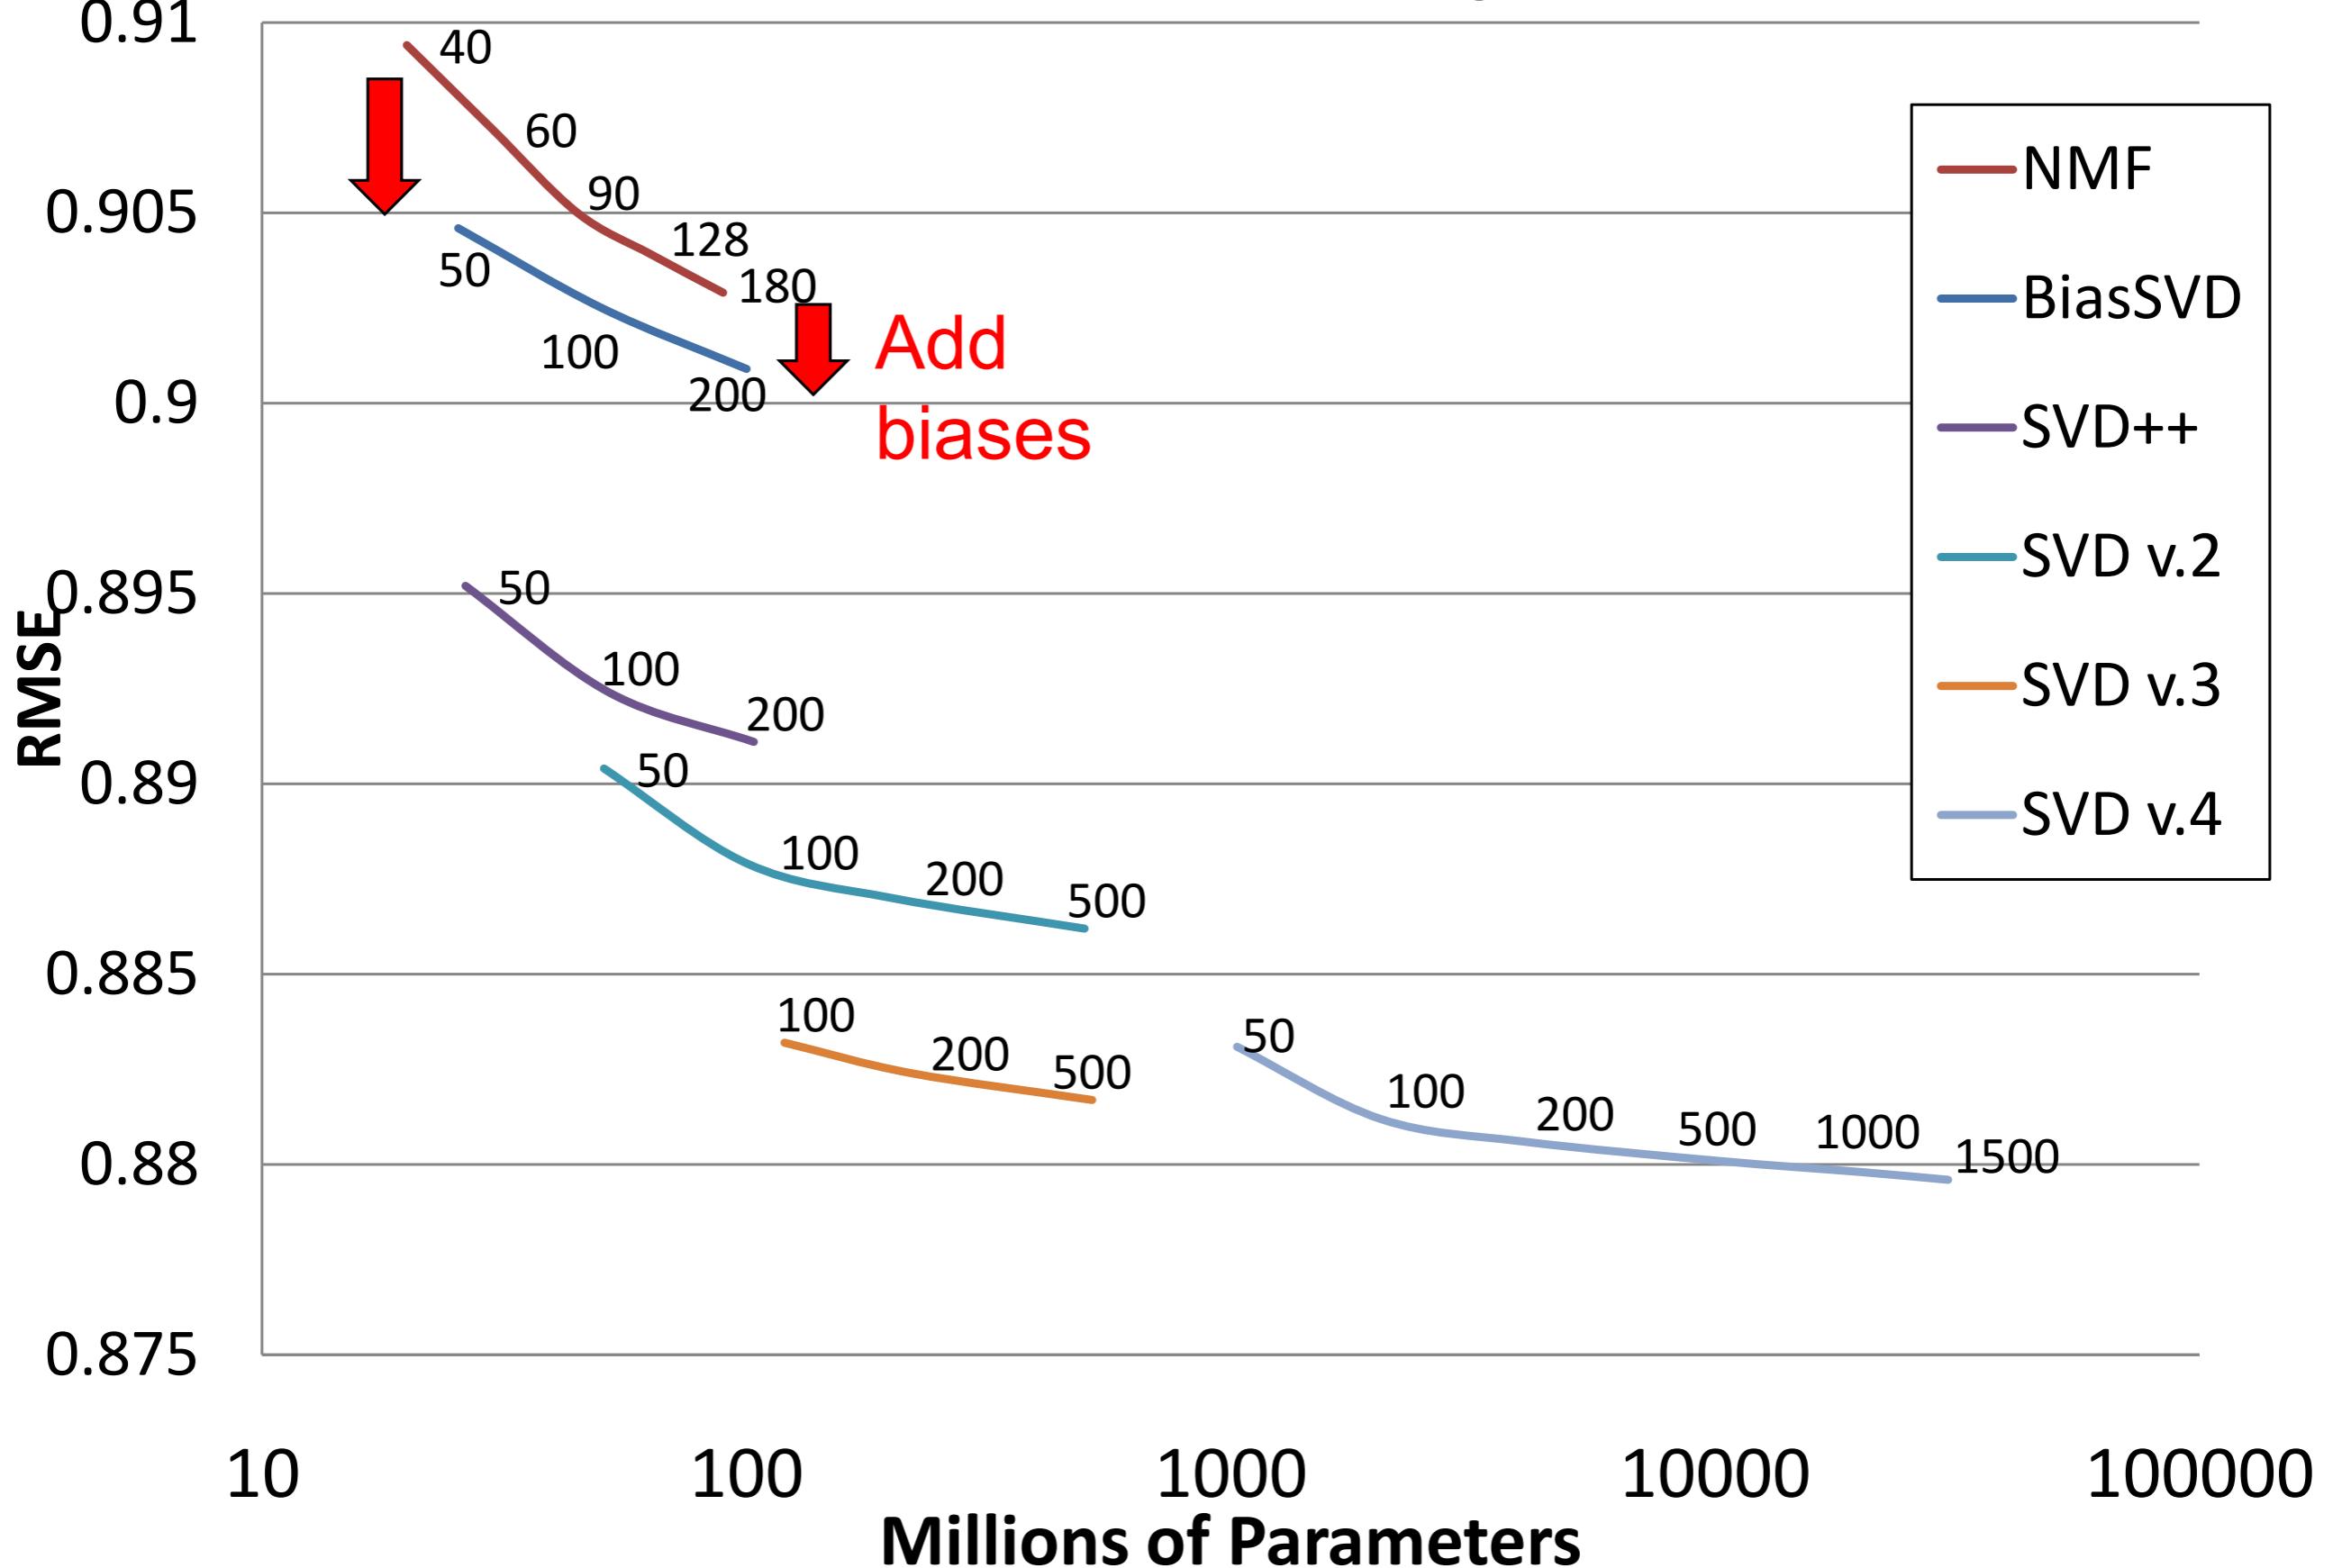
$$q_i \leftarrow (1 - \lambda \eta_t) q_i - \eta_t p_u (r_{ui} - \langle p_u, q_i \rangle)$$

faster

- No need for locking
- Multicore updates asynchronously
(Recht, Re, Wright, 2012 - Hogwild)

Improvements

Factor models: Error vs. #parameters



Bias

- **Objective Function**

$$\begin{aligned} \text{minimize}_{p,q} \quad & \sum_{(u,i) \in S} (r_{ui} - (\mu + b_u + b_i + \langle p_u, q_i \rangle))^2 + \\ & \lambda \left[\|p\|_{\text{Frob}}^2 + \|q\|_{\text{Frob}}^2 + \|b_{\text{users}}\|^2 + \|b_{\text{items}}\|^2 \right] \end{aligned}$$

- **Stochastic gradient descent**

$$p_u \leftarrow (1 - \lambda \eta_t) p_u - \eta_t q_i \rho_{ui}$$

$$q_i \leftarrow (1 - \lambda \eta_t) q_i - \eta_t p_u \rho_{ui}$$

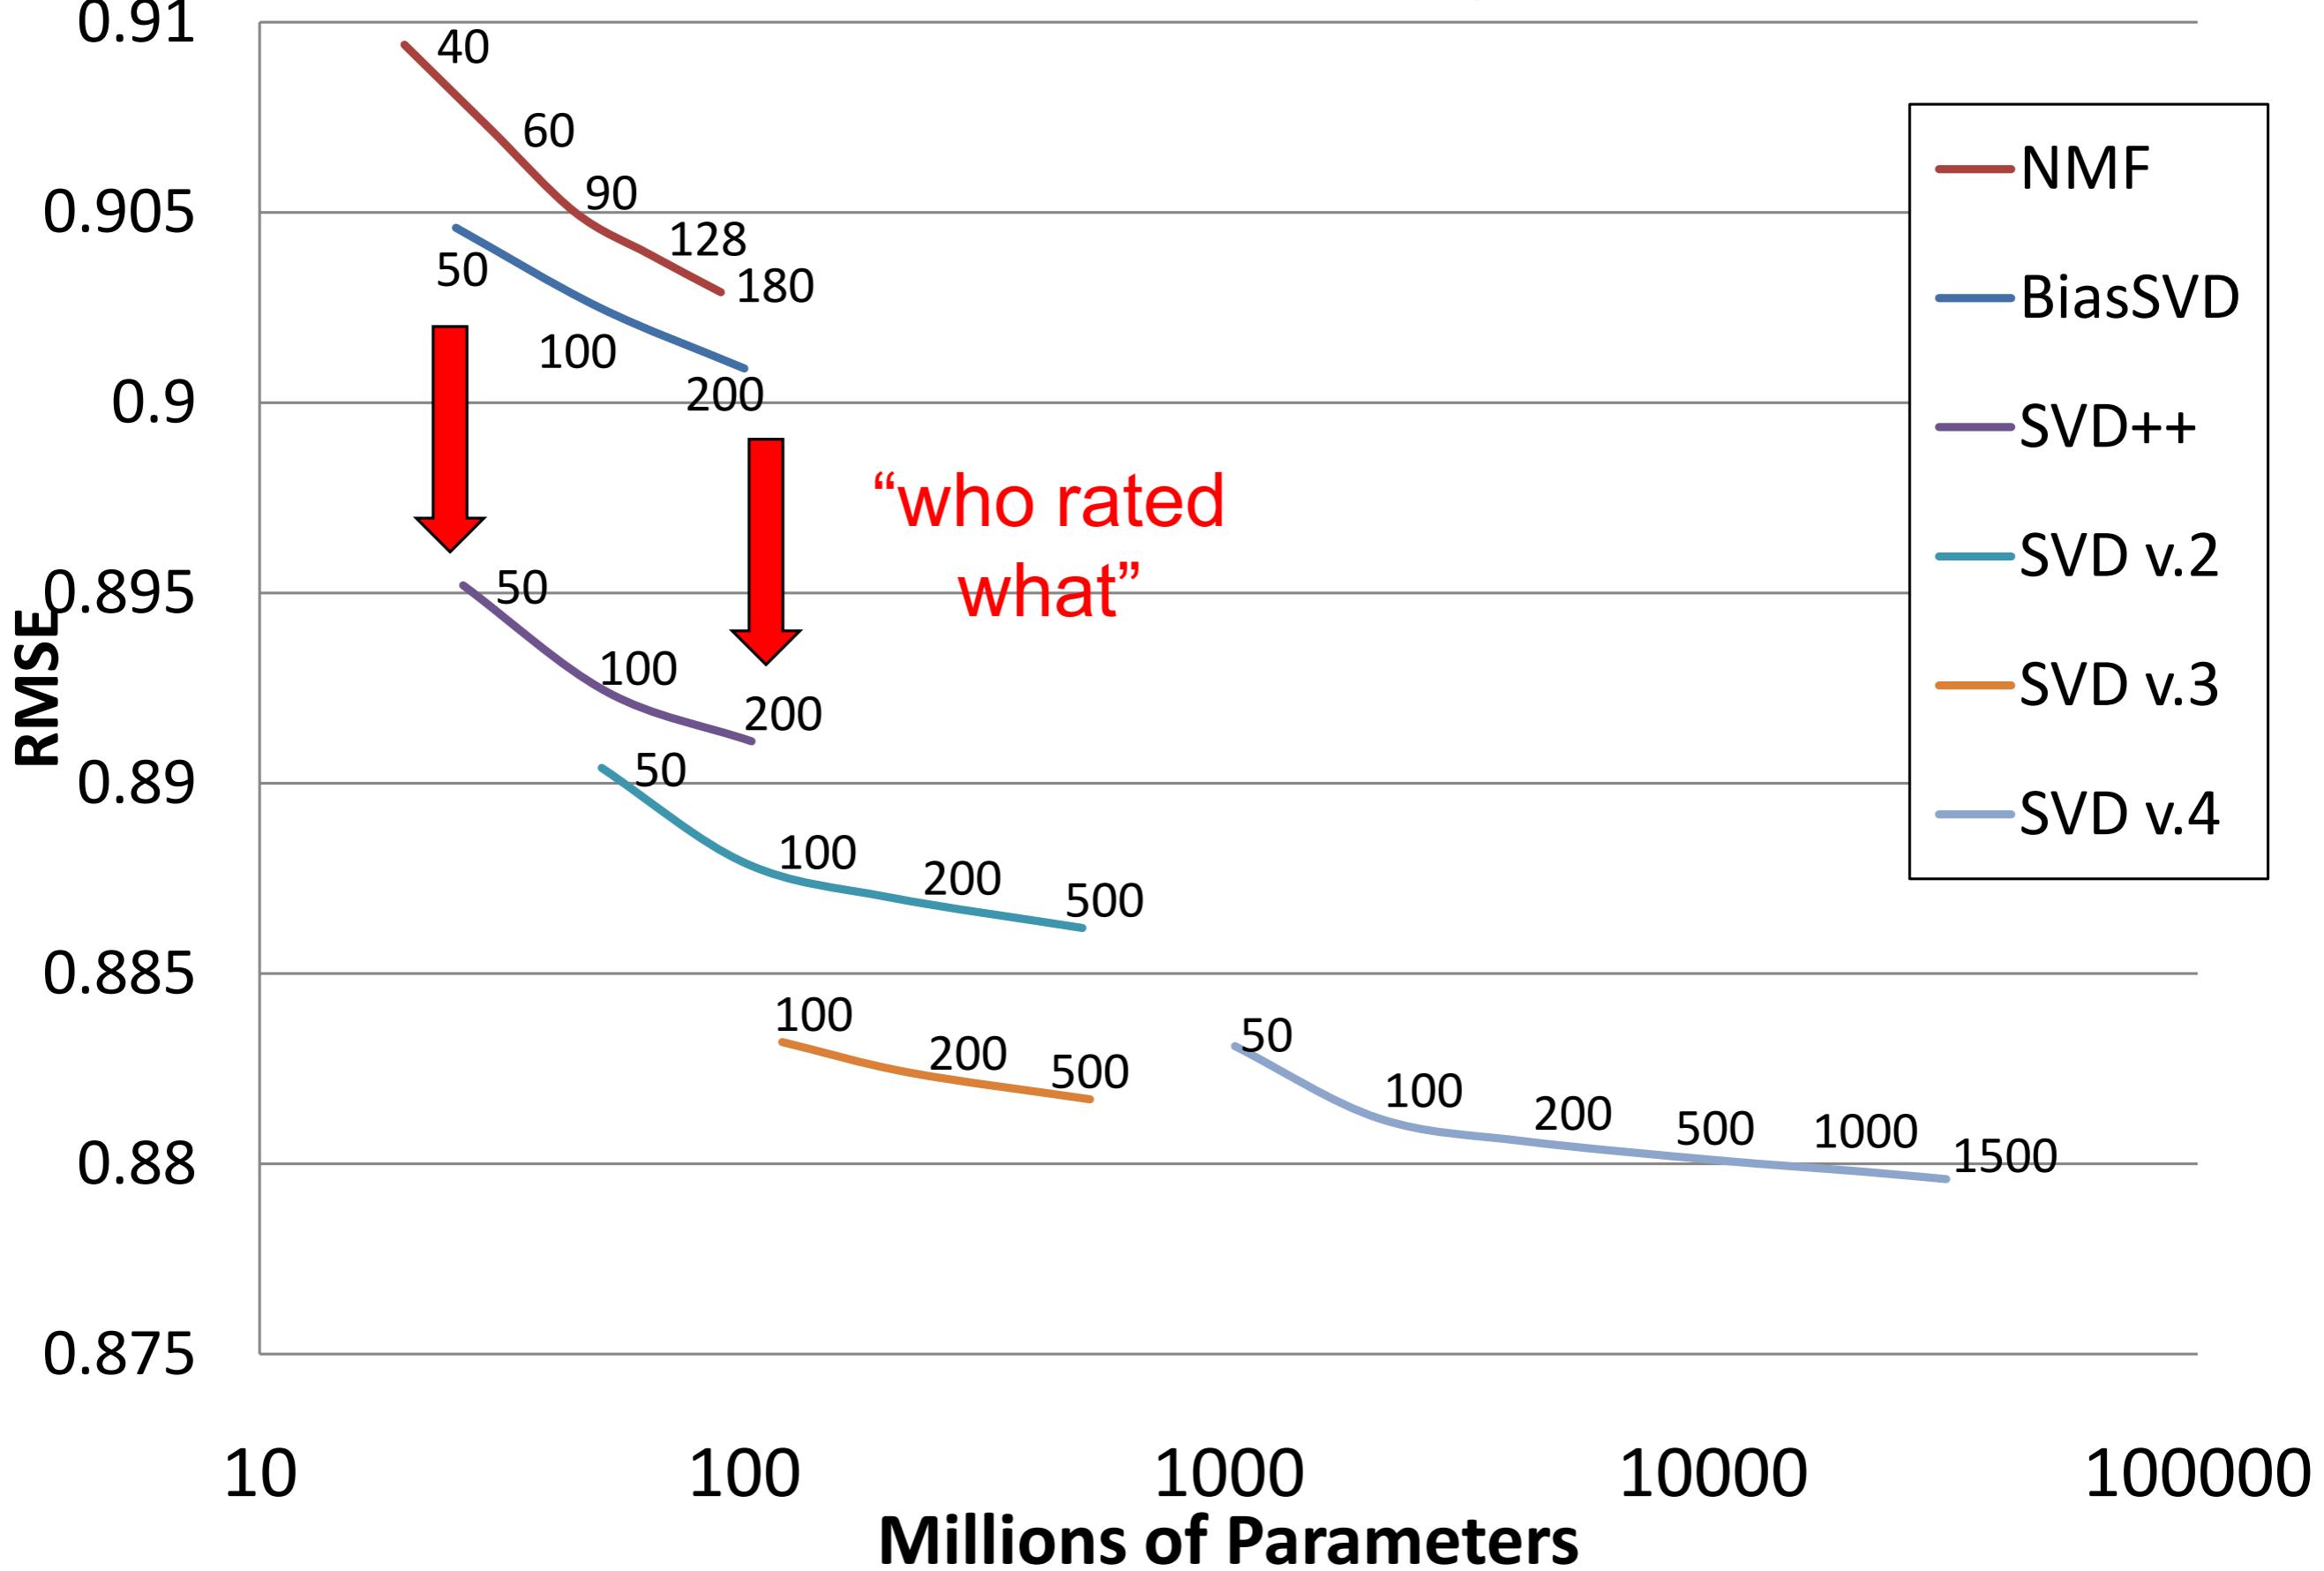
$$b_u \leftarrow (1 - \lambda \eta_t) b_u - \eta_t \rho_{ui}$$

$$b_i \leftarrow (1 - \lambda \eta_t) b_i - \eta_t \rho_{ui}$$

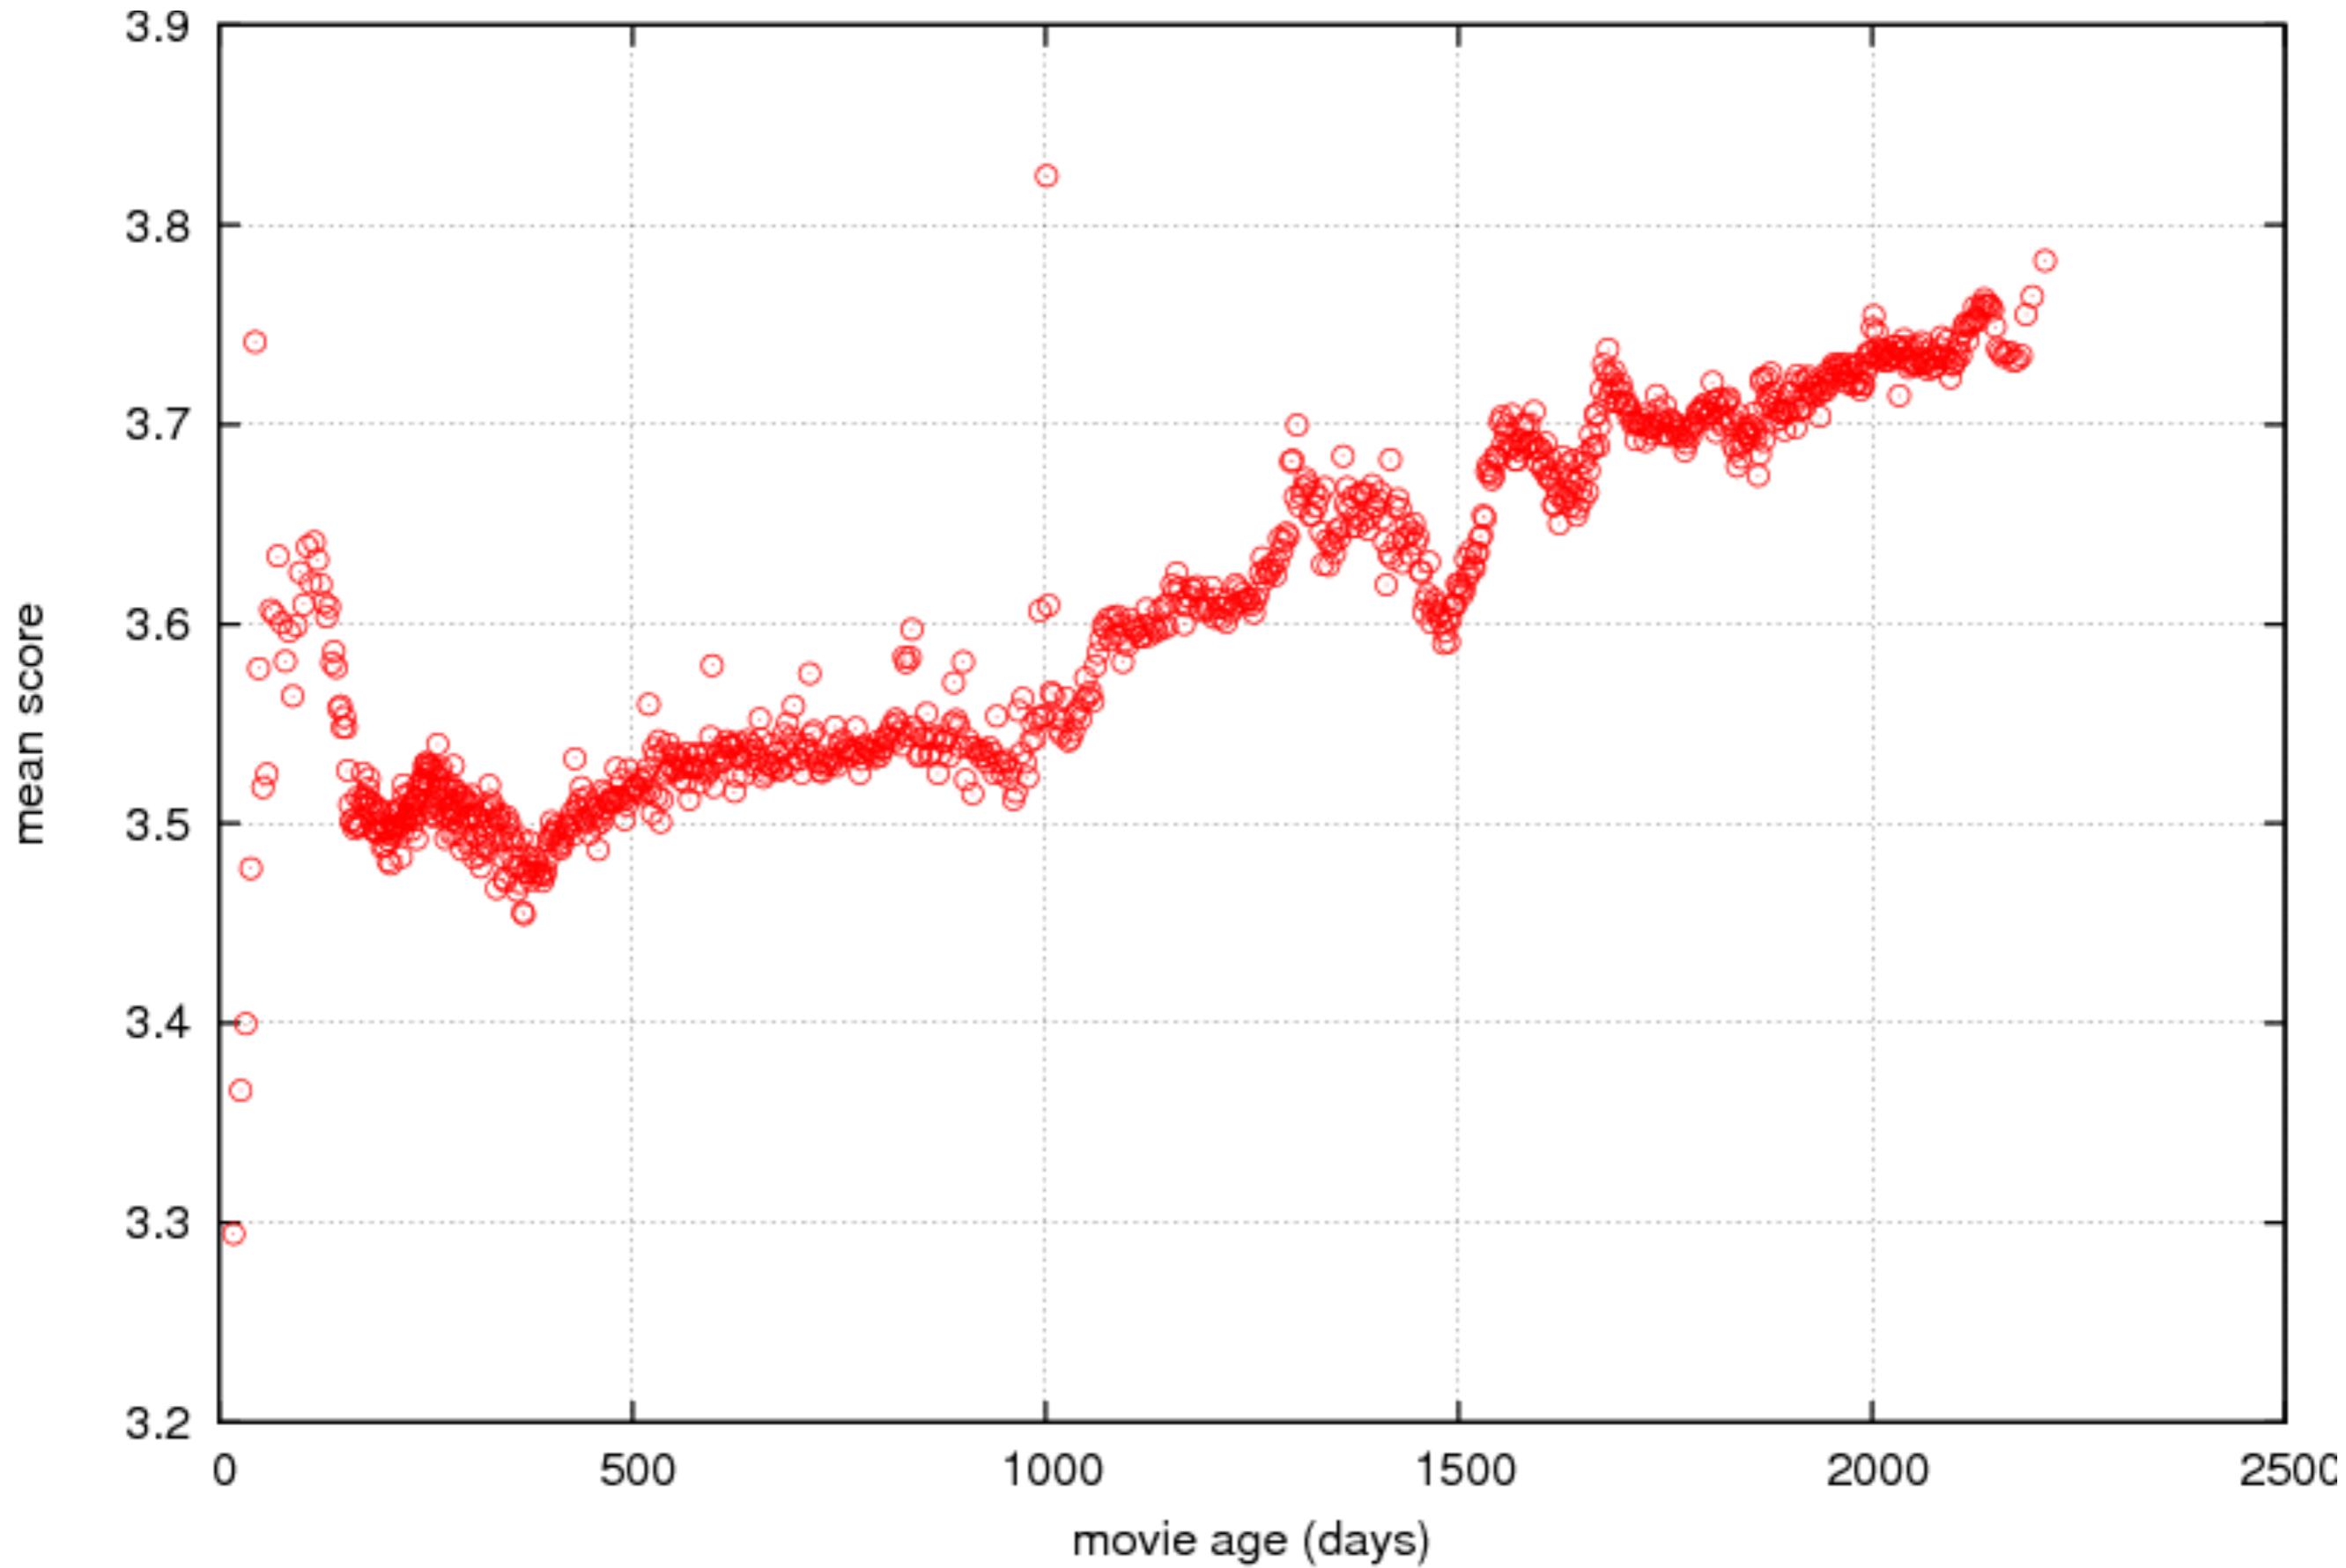
$$\mu \leftarrow (1 - \lambda \eta_t) \mu - \eta_t \rho_{ui}$$

where $\rho_{ui} = (r_{ui} - (\mu + b_i + b_u + \langle p_u, q_i \rangle))$

Factor models: Error vs. #parameters



Are movies getting better with time?



Sources of temporal change

- Items
 - Seasonal effects
(Christmas, Valentine's day, Holiday movies)
 - Public perception of movies (Oscar etc.)
- Users
 - Changed labeling of reviews
 - Anchoring (relative to previous movie)
 - Change of rater in household
 - Selection bias for time of viewing

Modeling temporal change

- Time-dependent bias
- Time-dependent user preferences

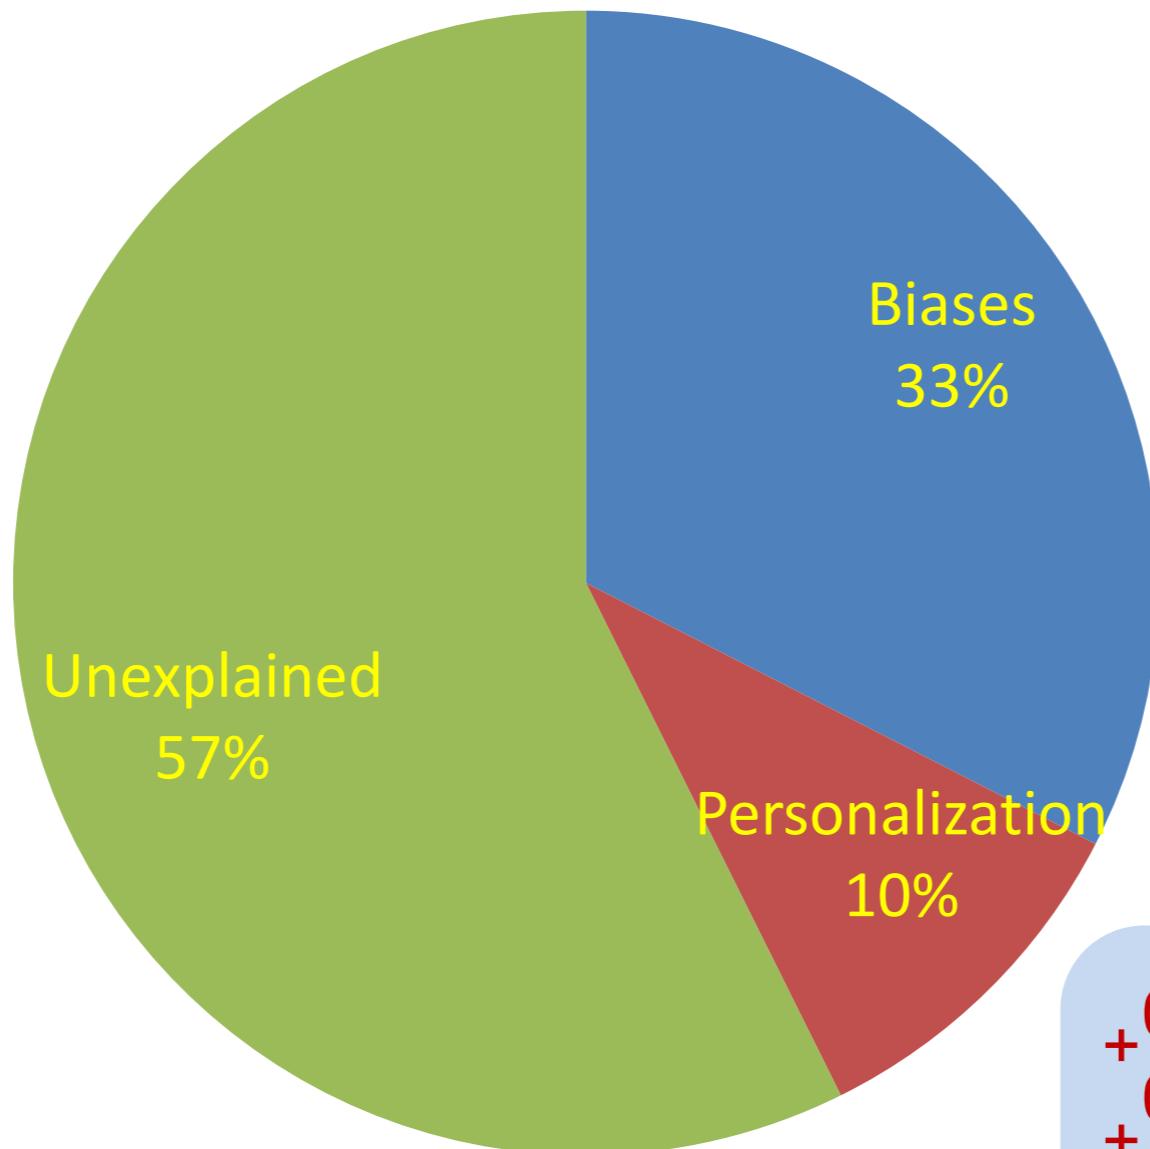
$$r_{ui}(t) = \mu + b_u(t) + b_i(t) + \langle q_i, p_u(t) \rangle$$

- Parameterize functions b and p
 - Slow changes for items
 - Fast sudden changes for users
 - Good parametrization is key

Koren et al., KDD 2009 (CF with temporal dynamics)

Bias matters

Sources of Variance in Netflix data



$$\begin{aligned} &+ 0.732 \text{ (unexplained)} \\ &+ 0.415 \text{ (biases)} \\ &+ 0.129 \text{ (personalization)} \\ \hline &1.276 \text{ (total variance)} \end{aligned}$$

Summary

- Neighborhood methods
 - User / movie similarity
 - Iteration on graph
- Matrix Factorization
 - Singular value decomposition
 - Convex reformulation
- Improvements

Further reading

- Neighborhood factorization
<http://research.yahoo.com/files/paper.pdf>
- Matrix Factorization for recommender systems
[https://datajobs.com/data-science-repo/
Recommender-Systems-%5BNetflix%5D.pdf](https://datajobs.com/data-science-repo/Recommender-Systems-%5BNetflix%5D.pdf)
- Yehuda Koren's papers
http://research.yahoo.com/Yehuda_Koren